

****Telecom Customer Retention Project**

Will Byrd, May 2024

Introduction**

This is a standard classification project. In this notebook, I will create classification models to predict whether or not a telephone service customer churned. Churn refers to customers who have canceled their subscriptions. I will focus on 2 different models-**Decision Tree Classifier and Logistic Regression**. A baseline of each model will be created and then iteratively improved upon until the best model is created. Methods to improve upon these models include

- feature scaling
- Oversampling(SMOTE)
- Pruning the Decision Tree
- Hyper Parameter Tuning

Data

The data used in this project is from Kaggle's Churn in Telecom's Dataset. This data is remarkably clean with no missing values and will allow me to focus on the principles of model building. Each record in this dataset represents a customer in Telecom and has attributes such as:



- state
- length of subscription
- type of plan
- usage
- whether or not the churned

We will be targeting churn values for customers. The churn column is our target column- otherwise known as our dependent variable. Accurately predicted customers who Churn will have values of True Positive and accurately predicted customers who do not Churn will have values of True Negative.

There will be a few preprocessing techniques on display here. We will notice some 1:1 correlations between our usage columns that will allow us to drop specific columns to simplify the data. We will also dummy code categorical variables such as International Plan and Voicemail Plan. Our state column will make building our models impossible, so we will use One Hot Encoding as well.

Goals

The main goal is to build the best model that can predict whether or not a customer will Churn. We want to be able to determine qualities in customers who will churn vs those who will not churn so the business can be more strategic and efficient. This can help the business allocate

resources either to markets that are more advantageous or to know how to better anticipate their budget. If a company can precisely predict how many customers will churn every month, they can be better prepared for the future.

Since we have an overwhelming majority of False values for Churn, we want to build a model that can most correctly predict true positives as this will give us confidence in knowing how much money (subscriptions) the company will lose month over month.

Overview

We will explore the data to better understand all features, correlations, and some distributions. Then build a baseline Decision Tree Model and refine it with the methods mentioned above. Every iteration of our baseline model will be evaluated with a cross validation score and Area Under the Curve (AUC) and then finally a confusion matrix will validate which model is the best. Then the process will start over again with a Logistic Regression Model.

Exploratory Data Analysis

Let's load in every library we can think of.

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 %matplotlib inline
        5 import warnings
        6 import seaborn as sns
        7 warnings.simplefilter(action='ignore', category=FutureWarning)
        8 #from sklearn.linear_model import LinearRegression
        9 from sklearn.preprocessing import OneHotEncoder
       10 #from seaborn import load_dataset
       11 from sklearn.model_selection import train_test_split
       12 from sklearn import tree
       13 from sklearn.tree import DecisionTreeClassifier
       14 from sklearn.tree import plot_tree
       15 from sklearn.preprocessing import StandardScaler
       16 #from sklearn.compose import ColumnTransformer
       17 #from sklearn.pipeline import Pipeline
       18 from sklearn.metrics import accuracy_score, classification_report
       19 from sklearn.linear_model import LogisticRegression
       20 #from sklearn.utils.class_weight import compute_class_weight
       21 #from imblearn.over_sampling import RandomOverSampler
       22 #from sklearn.model_selection import GridSearchCV
       23 from sklearn.model_selection import cross_val_score
       24 from sklearn.metrics import confusion_matrix
       25 from sklearn.metrics import roc_curve, auc
```

Let's take a look at the data to get an understanding of the features.

```
In [2]: ▶ 1 # using pandas to read in data
          2 df = pd.read_csv ('Data/telecom.csv')
          3 df.head()
```

Out[2]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	.
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	.
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	.
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	.
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	.
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	.

5 rows × 21 columns



We can see the State column will need to be one hot encoded. The International Plan and Voice Mail Plan columns will need to be dummy coded. And we can probably drop some columns that will not impact our results.



Let's clean this up and check for any missing values.

```
In [3]: ▶ 1 df.columns = df.columns.str.replace(' ', '_')
```

In [4]:  1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account_length                       3333 non-null   int64
2   area_code                           3333 non-null   int64
3   phone_number                        3333 non-null   object
4   international_plan                  3333 non-null   object
5   voice_mail_plan                     3333 non-null   object
6   number_vmail_messages               3333 non-null   int64
7   total_day_minutes                   3333 non-null   float64
8   total_day_calls                     3333 non-null   int64
9   total_day_charge                    3333 non-null   float64
10  total_eve_minutes                   3333 non-null   float64
11  total_eve_calls                     3333 non-null   int64
12  total_eve_charge                    3333 non-null   float64
13  total_night_minutes                 3333 non-null   float64
14  total_night_calls                   3333 non-null   int64
15  total_night_charge                  3333 non-null   float64
16  total_intl_minutes                  3333 non-null   float64
17  total_intl_calls                    3333 non-null   int64
18  total_intl_charge                   3333 non-null   float64
19  customer_service_calls              3333 non-null   int64
20  churn                               3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```



Great, no Null values. Let's take a look at all values in all columns to get a better understanding.

In [5]:  1 *#for col in df.columns:*
2 *# print(f"{col} vals: {df[col].unique()} \n")*

Everything looks normal so far, except the area codes. Only 3 area codes makes me think that column is not going to be usable. We can also assume phone number has no bearing on whether a customer churns. We have a few categorical variables we can dummy code:

- intl plan
- voicemail plan

```
In [6]: 1 df.replace({'no': 0, 'yes':1, 'false':0, 'true':1}, inplace=True)
        2 df.head()
        3
```

Out[6]:

	state	account_length	area_code	phone_number	international_plan	voice_mail_plan	nu
0	KS	128	415	382-4657	0	1	
1	OH	107	415	371-7191	0	1	
2	NJ	137	415	358-1921	0	0	
3	OH	84	408	375-9999	1	0	
4	OK	75	415	330-6626	1	0	

5 rows × 21 columns

```
In [7]: 1 df.drop('phone_number', axis = 1, inplace = True)
```

```
In [8]: 1 df.drop('area_code', axis = 1, inplace = True)
```

I'm thinking minutes, charges, and calls are all related. Let's take a deeper dive into this and figure it out.



```
In [9]: 1 mins_calls_charge = df[['total_day_minutes', 'total_day_calls', 'total_eve_minutes', 'total_eve_calls', 'total_day_charge', 'total_eve_charge']]
        2 mins_calls_charge.head()
```

Out[9]:

	total_day_minutes	total_day_calls	total_day_charge	total_eve_minutes	total_eve_calls	total_eve_charge
0	265.1	110	45.07	197.4	99	35.49
1	161.6	123	27.47	195.5	103	35.49
2	243.4	114	41.38	121.2	110	20.91
3	299.4	71	50.90	61.9	88	10.90
4	166.7	113	28.34	148.3	122	20.91

In [10]: 1 mins_calls_charge.corr()

Out[10]:

	total_day_minutes	total_day_calls	total_day_charge	total_eve_minutes
total_day_minutes	1.000000	0.006750	1.000000	0.007043
total_day_calls	0.006750	1.000000	0.006753	-0.021451
total_day_charge	1.000000	0.006753	1.000000	0.007050
total_eve_minutes	0.007043	-0.021451	0.007050	1.000000
total_eve_calls	0.015769	0.006462	0.015769	-0.011430
total_eve_charge	0.007029	-0.021449	0.007036	1.000000
total_night_minutes	0.004323	0.022938	0.004324	-0.012584
total_night_calls	0.022972	-0.019557	0.022972	0.007586
total_night_charge	0.004300	0.022927	0.004301	-0.012593
total_intl_minutes	-0.010155	0.021565	-0.010157	-0.011035
total_intl_calls	0.008033	0.004574	0.008032	0.002541
total_intl_charge	-0.010092	0.021666	-0.010094	-0.011067

OK cool. We have direct 1:1 correlations between minutes and charges. For that reason, we can drop the minutes columns.



In [11]: 1 df.drop('total_day_minutes', axis = 1, inplace = True)

In [12]: 1 df.drop('total_eve_minutes', axis = 1, inplace = True)

In [13]: 1 df.drop('total_night_minutes', axis = 1, inplace = True)

In [14]: 1 df.drop('total_intl_minutes', axis = 1, inplace = True)

In [15]: 1 df.head()

Out[15]:

	state	account_length	international_plan	voice_mail_plan	number_vmail_messages	total
0	KS	128	0	1		25
1	OH	107	0	1		26
2	NJ	137	0	0		0
3	OH	84	1	0		0
4	OK	75	1	0		0

```
In [16]: 1 df['churn'].value_counts()
```

```
Out[16]: False    2850
         True     483
         Name: churn, dtype: int64
```

Let's establish our X and Y variables. This will allow us to build our test and train sets for modeling.

Our X values are independent variables. In this dataset that includes all values except Churn.

```
In [17]: 1 X = df.drop('churn', axis = 1)
         2 y = df.churn
```

Now that our X and Y values are established, we can prepare our training sets for modeling.

```
In [18]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
```

One Hot Encoding

Let's one hot encode our state column so we can build our models. This is always good practice for categorical variables that aren't binary. This will create extra columns and we will have to drop our original State column as a result.



```
In [19]: 1 # instantiate ohe object
         2 ohe = OneHotEncoder(sparse = False, handle_unknown = "ignore")
         3
         4 # fit ohe on small train data
         5 ohe.fit(X_train[['state']])
         6
         7 # access the column names of the states
         8 col_names = ohe.categories_[0]
         9
        10 # make a df with encoded states
        11 train_state_encoded = pd.DataFrame(ohe.transform(X_train[["state"]]),
        12                                   index = X_train.index,
        13                                   columns = col_names)
        14
        15 # combine encoded states with X_train and drop old 'state' column
        16 X_train = pd.concat([X_train.drop("state", axis = 1), train_state_encoded])
```

```
In [20]: 1 X_train.head()
```

Out[20]:

	account_length	international_plan	voice_mail_plan	number_vmail_messages	total_d
367	45	0	0	0	
3103	115	0	0	0	
549	121	0	1	31	
2531	180	0	0	0	
2378	112	0	0	0	

5 rows × 64 columns

Let's inspect our training set below. This is a portion of our original df, so expect it to be slightly smaller than the original.



In [21]: ▶ 1 X_train.info()



```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 2499 entries, 367 to 3174
```

```
Data columns (total 64 columns):
```

#	Column	Non-Null	Count	Dtype
0	account_length	2499	non-null	int64
1	international_plan	2499	non-null	int64
2	voice_mail_plan	2499	non-null	int64
3	number_vmail_messages	2499	non-null	int64
4	total_day_calls	2499	non-null	int64
5	total_day_charge	2499	non-null	float64
6	total_eve_calls	2499	non-null	int64
7	total_eve_charge	2499	non-null	float64
8	total_night_calls	2499	non-null	int64
9	total_night_charge	2499	non-null	float64
10	total_intl_calls	2499	non-null	int64
11	total_intl_charge	2499	non-null	float64
12	customer_service_calls	2499	non-null	int64
13	AK	2499	non-null	float64
14	AL	2499	non-null	float64
15	AR	2499	non-null	float64
16	AZ	2499	non-null	float64
17	CA	2499	non-null	float64
18	CO	2499	non-null	float64
19	CT	2499	non-null	float64
20	DC	2499	non-null	float64
21	DE	2499	non-null	float64
22	FL	2499	non-null	float64
23	GA	2499	non-null	float64
24	HI	2499	non-null	float64
25	IA	2499	non-null	float64
26	ID	2499	non-null	float64
27	IL	2499	non-null	float64
28	IN	2499	non-null	float64
29	KS	2499	non-null	float64
30	KY	2499	non-null	float64
31	LA	2499	non-null	float64
32	MA	2499	non-null	float64
33	MD	2499	non-null	float64
34	ME	2499	non-null	float64
35	MI	2499	non-null	float64
36	MN	2499	non-null	float64
37	MO	2499	non-null	float64
38	MS	2499	non-null	float64
39	MT	2499	non-null	float64
40	NC	2499	non-null	float64
41	ND	2499	non-null	float64
42	NE	2499	non-null	float64
43	NH	2499	non-null	float64
44	NJ	2499	non-null	float64
45	NM	2499	non-null	float64
46	NV	2499	non-null	float64
47	NY	2499	non-null	float64
48	OH	2499	non-null	float64
49	OK	2499	non-null	float64
50	OR	2499	non-null	float64
51	PA	2499	non-null	float64



```

52  RI      2499 non-null    float64
53  SC      2499 non-null    float64
54  SD      2499 non-null    float64
55  TN      2499 non-null    float64
56  TX      2499 non-null    float64
57  UT      2499 non-null    float64
58  VA      2499 non-null    float64
59  VT      2499 non-null    float64
60  WA      2499 non-null    float64
61  WI      2499 non-null    float64
62  WV      2499 non-null    float64
63  WY      2499 non-null    float64
dtypes: float64(55), int64(9)
memory usage: 1.2 MB

```

Ok, our data is cleaned up and split into test and train sets. We can begin to build some models. First, I want some more info on our categorical features.

```
In [22]: 1 df.churn.value_counts()
```

```
Out[22]: False    2850
         True     483
         Name: churn, dtype: int64
```

```
In [23]: 1 df.churn.value_counts()/len(df.churn)
```

```
Out[23]: False    0.855086
         True     0.144914
         Name: churn, dtype: float64
```

<15% of customers churn. This is actually industry standard, so all good. here.

```
In [24]: 1 df.international_plan.value_counts()
```

```
Out[24]: 0    3010
         1     323
         Name: international_plan, dtype: int64
```

```
In [25]: 1 df.international_plan.value_counts()/len(df.international_plan)
```

```
Out[25]: 0    0.90309
         1    0.09691
         Name: international_plan, dtype: float64
```

<10% of customers have international plans.

```
In [26]: 1 df.voice_mail_plan.value_counts()/len(df.voice_mail_plan)
```

```
Out[26]: 0    0.723372
         1    0.276628
         Name: voice_mail_plan, dtype: float64
```

<30% have voice mail plans.

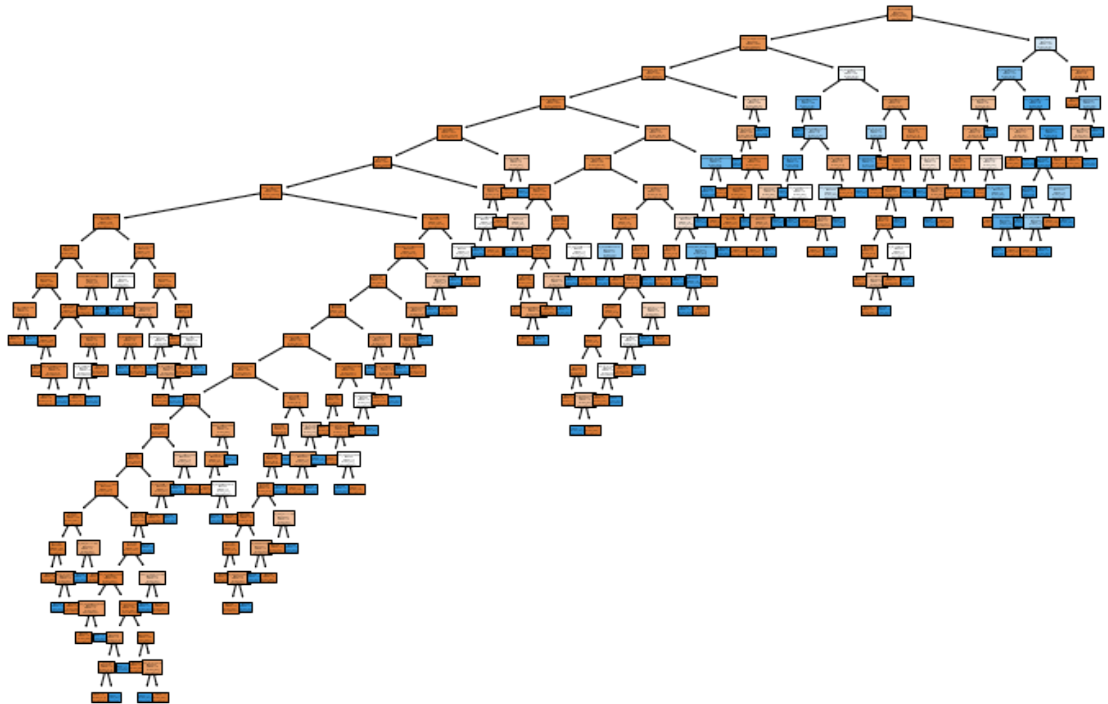
Baseline Decision Tree Model

Let's build our baseline model. All subsequent decision tree models will be evaluated against this baseline.

```
In [27]: 1 # One-hot encode categorical features
2 X_train_encoded = pd.get_dummies(X_train)
3 X_test_encoded = pd.get_dummies(X_test)

In [28]: 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model = DecisionTreeClassifier().fit(X_train_encoded, y_train)
3
4 # Predict the labels for the test set
5 y_pred_dt = decision_tree_model.predict(X_test_encoded)
6
7 # Now, you can proceed with plotting the decision tree or any other analysis
8 plt.figure(figsize=(12, 8))
9 plot_tree(decision_tree_model, filled=True, feature_names=X_train_encoded.columns)
10 plt.title("Decision Tree Visualization")
11 plt.show()
12
13
```

Decision Tree Visualization



```
In [29]: ▶ 1 # Evaluate the model
2 accuracy_dt = accuracy_score(y_test, y_pred_dt)
3 classification_rep_dt = classification_report(y_test, y_pred_dt)
4
5 print("Decision Tree Model Evaluation:")
6 print("Accuracy:", accuracy_dt)
7 print("Classification Report:")
8 print(classification_rep_dt)
9
```

Decision Tree Model Evaluation:

Accuracy: 0.9112709832134293

Classification Report:

	precision	recall	f1-score	support
False	0.95	0.94	0.95	709
True	0.69	0.74	0.71	125
accuracy			0.91	834
macro avg	0.82	0.84	0.83	834
weighted avg	0.91	0.91	0.91	834

This is actually a really strong baseline model. However, we want to make sure we are focusing on the right metrics, so we will add a cross validation score and AUC score as well.



To add more here:

Precision is the number of True Positives/all predicted positives or the True Positive Rate (TPR). So Precision in this case is how often the model correctly predicts the target class.

Recall is the number of True Positives/actual positives. In this case, Recall is the models ability to find all objects of the target class.

Cross Validation

Cross Validation is a technique used to partition a dataset into multiple subsets for training. This will help detect overfitting and give us more confidence in our model.

```
In [30]: ▶ 1 cv_scores = cross_val_score(decision_tree_model, X_test_encoded, y_te
2 print('Cross-Validation Scores', cv_scores)
3 print('Mean CV Score', cv_scores.mean())
```

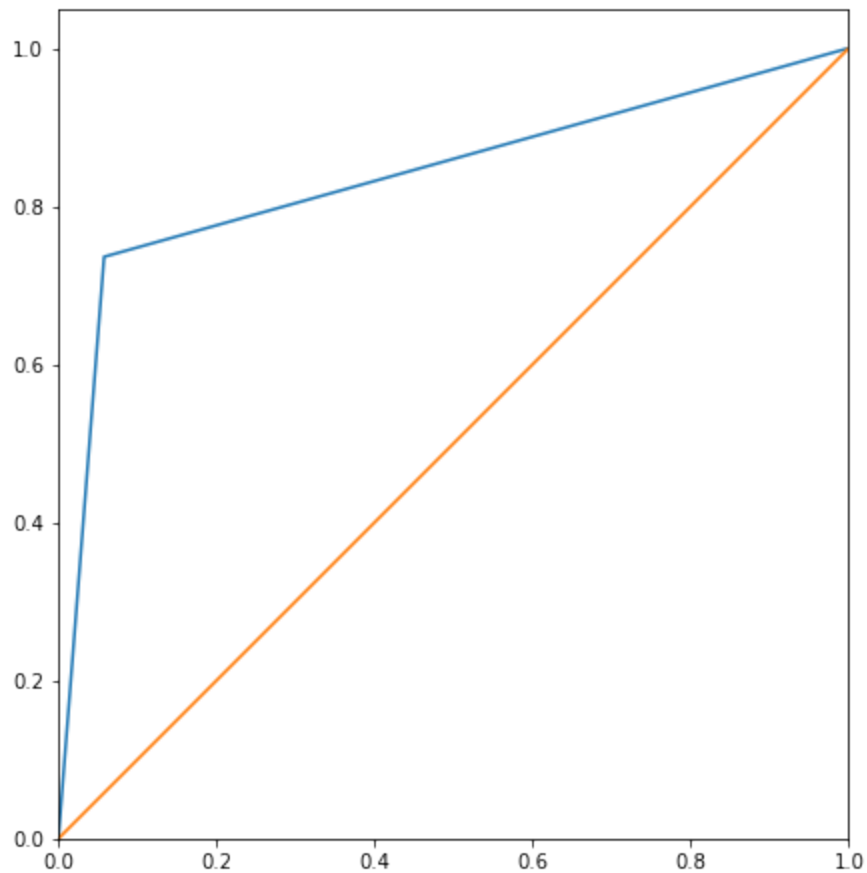
Cross-Validation Scores [0.91017964 0.92215569 0.89221557 0.91017964 0.87349398]

Mean CV Score 0.9016449029651541

Area Under the Curve

AUC measures the overall performance of our binary classification model. We can see our True Positive Rate (TPR) on the Y-axis and our False Positive Rate (FPR) on the X-axis. These values range from 0 to 1.

```
In [31]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_encoded = decision_tree_model.predict_proba(X_test_encoded)[:],
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_encoded)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
18
19
```



0.839086036671368



In [32]: 1 `print(y_test)`

```
438      False
2674     False
1345      True
1957     False
2148     False
...
3257     False
1586     False
3068     False
2484     False
219      False
Name: churn, Length: 834, dtype: bool
```

We have made our baseline model and evaluated it on a variety of metrics. now let's take a quick look at Feature importance. This will help us understand which features (independent variables) are impacting our dependent variable the most.

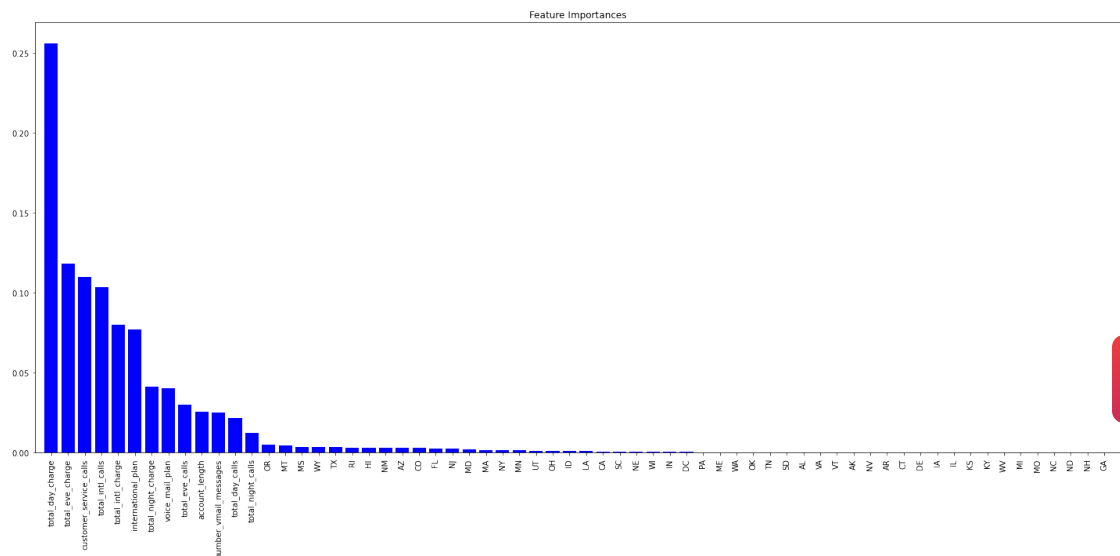
Feature Importance



```

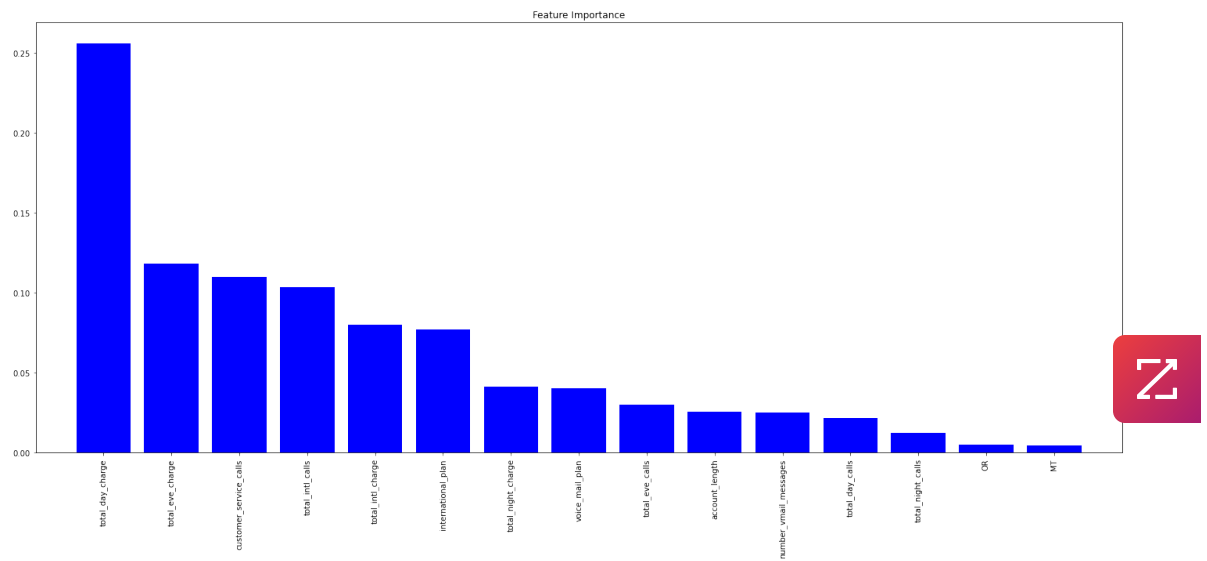
In [33]: ▶ 1 X_train_final= pd.DataFrame(X_train_encoded,columns=X_train.columns)
2
3 feature_importances = decision_tree_model.feature_importances_
4
5 feature_names = X_train_final.columns
6 indices = np.argsort(feature_importances)[::-1]
7
8 plt.figure(figsize=(20, 10))
9 plt.title("Feature Importances")
10 plt.bar(range(X_train_final.shape[1]), feature_importances[indices],
11 plt.xticks(range(X_train_final.shape[1]), feature_names[indices], rot
12 plt.xlim([-1, X_train_final.shape[1]])
13 plt.tight_layout()
14 plt.show()

```




Lots of unimportant features. Lets reduce this list down to the top 15, since most of the states don't carry much weight here.


```
In [34]: 1 X_train_final= pd.DataFrame(X_train_encoded,columns=X_train.columns)
2
3 feature_importances = decision_tree_model.feature_importances_
4
5 feature_names = X_train_final.columns
6 indices = np.argsort(feature_importances)[::-1][:15]
7
8 plt.figure(figsize=(20, 10))
9 plt.title("Feature Importance")
10 plt.bar(range(15), feature_importances[indices], color="b", align="ce
11 plt.xticks(range(15), feature_names[indices], rotation=90)
12 plt.xlim([-1, 15])
13 plt.tight_layout()
14 plt.show()
```



We can see that the most impactful features are:

- total_day_charge
- total_eve_charge
- customer_service_calls

In [35]:  1 df.corr().churn.sort_values(ascending=False)

Out[35]:


churn	1.000000
international_plan	0.259852
customer_service_calls	0.208750
total_day_charge	0.205151
total_eve_charge	0.092786
total_intl_charge	0.068259
total_night_charge	0.035496
total_day_calls	0.018459
account_length	0.016541
total_eve_calls	0.009233
total_night_calls	0.006141
total_intl_calls	-0.052844
number_vmail_messages	-0.089728
voice_mail_plan	-0.102148

Name: churn, dtype: float64

hmm, this is slightly different than what we are seeing above. However, we can deal with this later.

Let's take a look at the **Categorical Features** above and check the ratios.

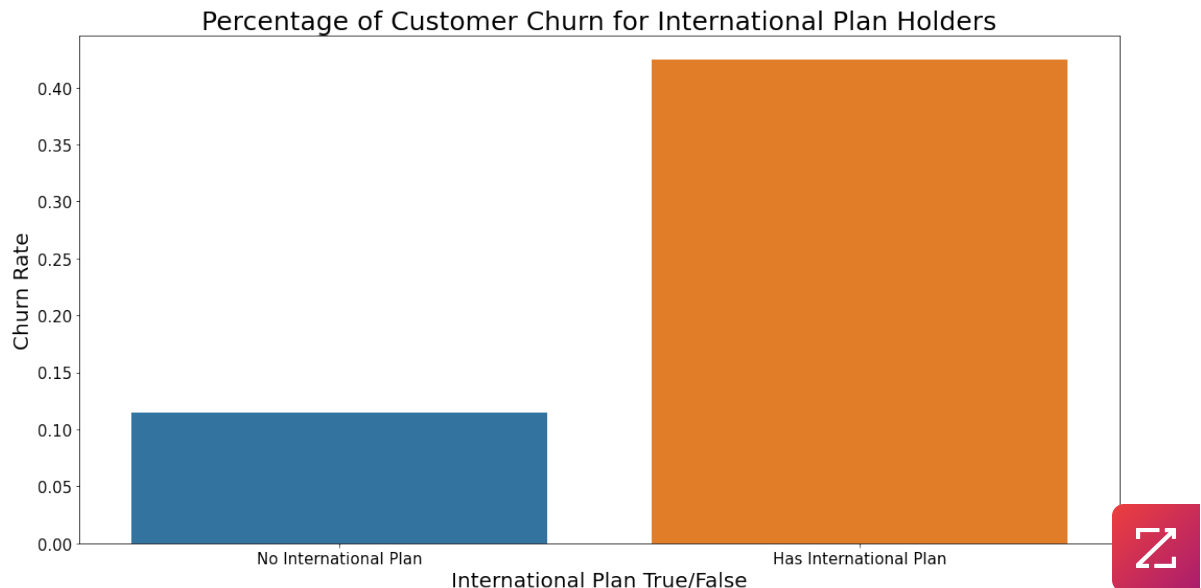
1st up is the International Plan.

In [36]:  1 int_plan_churn = pd.DataFrame(df.groupby(['international_plan'])['churn'].mean())
2 int_plan_churn

Out[36]:

churn	
international_plan	
0	0.114950
1	0.424149

```
In [37]: 1 fig, ax = plt.subplots(figsize=(15,8))
2 sns.barplot(x = [0, 1], y = 'churn', data = int_plan_churn, ax = ax)
3 plt.title('Percentage of Customer Churn for International Plan Holder
4 ax.tick_params(axis = 'both', labelsiz = 15)
5 plt.xlabel('International Plan True/False', fontsize = 20)
6 plt.ylabel('Churn Rate', fontsize = 20)
7 ax.set_xticklabels(['No International Plan', 'Has International Plan']
8 plt.tight_layout()
9
```



OK! we can see that customers who have international plans churn at a much higher rate than customers who don't. Maybe they are unhappy with their monthly bill? Maybe taking a look at customer service calls could shed some light on this. I would assume customers who make more customer service calls are probably not happy customers and therefore churning.

Let's take a look at the values for **Customer Service Calls** next.

```
In [38]: 1 df['customer_service_calls'].value_counts()
```

```
Out[38]: 1 1181
2 759
0 697
3 429
4 166
5 66
6 22
7 9
9 2
8 2
Name: customer_service_calls, dtype: int64
```

Most customers aren't making that many customer service calls. I'm assuming everyone makes one setting up their phone plan. Let's look for some **Correlations**.

```
In [39]: ▶ 1 csc = pd.DataFrame(df.groupby(['customer_service_calls'])['churn'].me  
2 csc
```

Out[39]:

		churn
customer_service_calls		
0		0.131994
1		0.103302
2		0.114625
3		0.102564
4		0.457831
5		0.606061
6		0.636364
7		0.555556
8		0.500000
9		1.000000

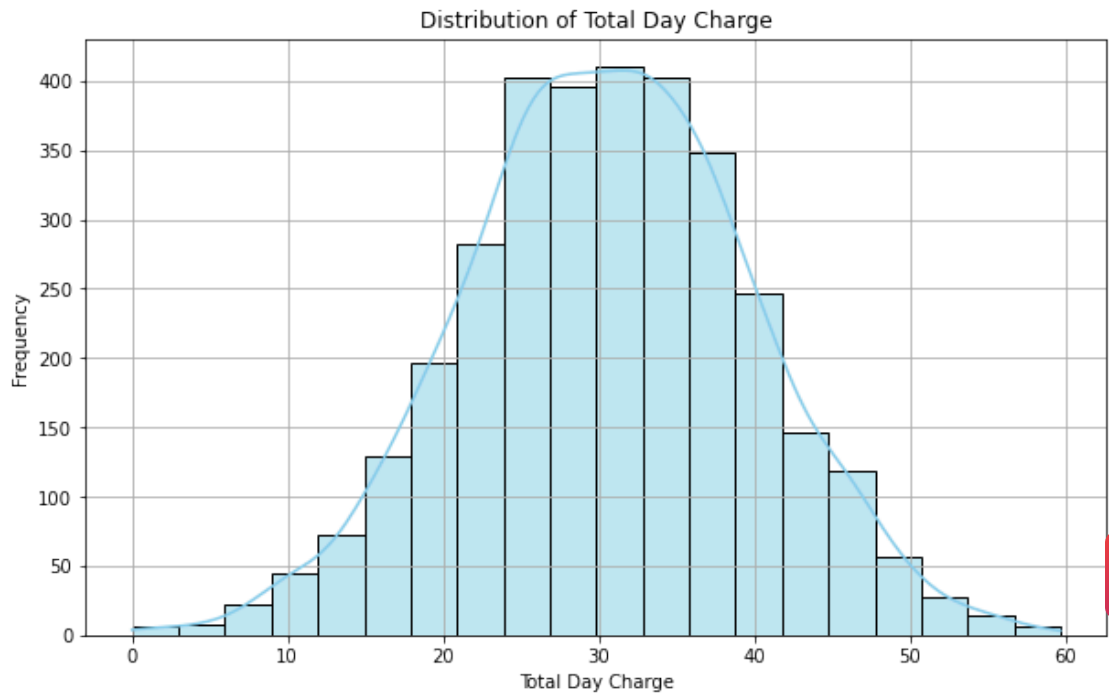
As we suspected, customers who make more customer service calls tend to churn at higher rates.

Maybe the company needs to look into better customer service.

Let's take a look at the Distribution of **Total Day Charge**.



```
In [40]: ▶ 1 total_day_charge_values = df['total_day_charge']
2 plt.figure(figsize=(10, 6))
3 sns.histplot(total_day_charge_values, bins=20, color='skyblue', edgecolor='black')
4 plt.title('Distribution of Total Day Charge')
5 plt.xlabel('Total Day Charge')
6 plt.ylabel('Frequency')
7 plt.grid(True)
8 plt.show()
```



It's a **normal distribution**. I'm going to find the mean of 'total_day_charge' to see if there is a threshold for customers who churn once they spend a certain amount.

```
In [41]: 1 tdc = pd.DataFrame(df.groupby(['total_day_charge'])['churn'].mean())
        2 tdc
```

Out[41]:

	churn
total_day_charge	
0.00	0.5
0.44	0.0
1.33	0.0
1.34	0.0
2.13	0.0
...	...
57.04	1.0
57.36	1.0
58.70	1.0
58.96	1.0

```
In [42]: 1 # Assuming 'total_day_charge' is a column in your DataFrame df
        2 total_day_charge_range = df['total_day_charge'].describe()
        3
        4 print("Range of 'total_day_charge' column:")
        5 print("Minimum:", total_day_charge_range['min'])
        6 print("Maximum:", total_day_charge_range['max'])
        7
```

```
Range of 'total_day_charge' column:
Minimum: 0.0
Maximum: 59.64
```

We can see the max 'total_day_charge' value is \$59.64/day. And we can see that as customers get close to that value, they have a 100% churn rate.

Improving The Baseline Model

We have done some evaluation of our categorical features and determined the following customers are **likely to churn**:

- customers who make multiple customer service calls
- customers who have a high daily bill
- customers who have an international plan

We can see that the same will be true for customers with high eve and night charges as well.

Now let's try to improve upon our baseline model. We can do the following to improve upon our baseline model:

- Feature Scaling

- Under/Over sampling
- Hyper Parameter Tuning

After we have improved this model, we will build a logistic regression model as well.

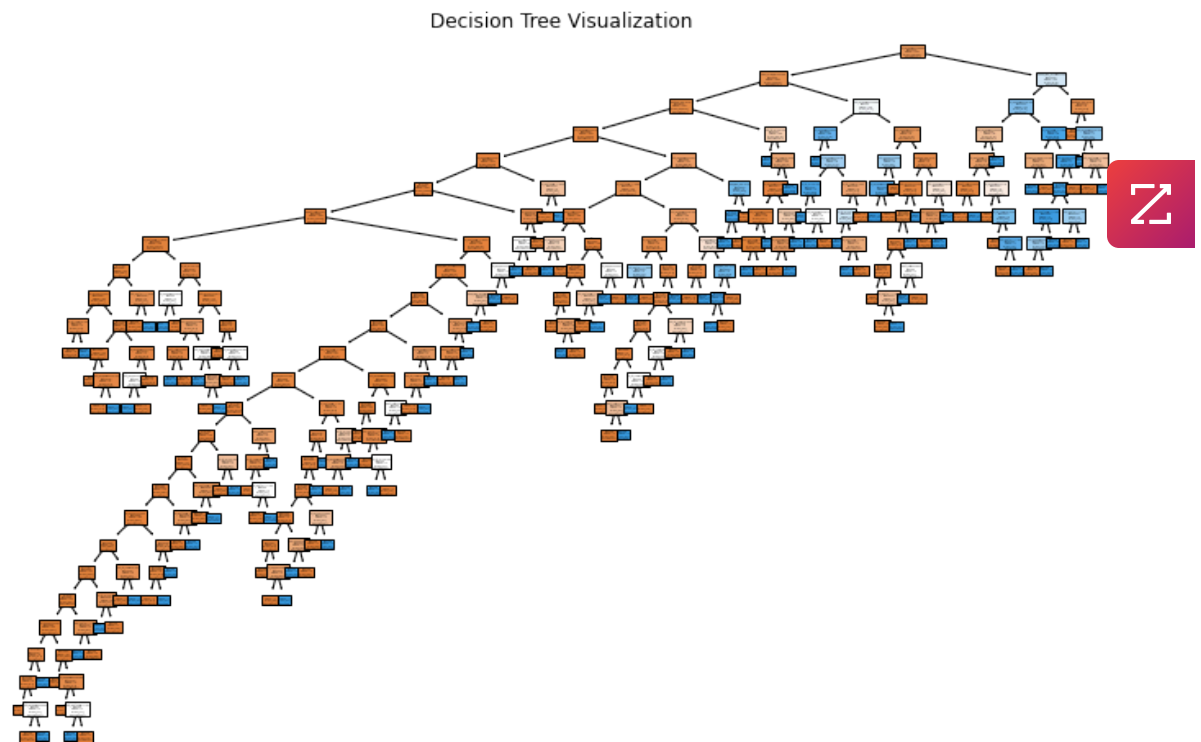
I'll move our Decision Tree Classifier back down here for reference:



```

In [43]: ▶ 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model = DecisionTreeClassifier().fit(X_train_encoded, y_train)
3
4 # Predict the labels for the test set
5 y_pred_dt = decision_tree_model.predict(X_test_encoded)
6
7 # Now, you can proceed with plotting the decision tree or any other analysis
8 plt.figure(figsize=(12, 8))
9 plot_tree(decision_tree_model, filled=True, feature_names=X_train_encoded.feature_names_)
10 plt.title("Decision Tree Visualization")
11 plt.show()
12 # Evaluate the model
13 accuracy_dt = accuracy_score(y_test, y_pred_dt)
14 classification_rep_dt = classification_report(y_test, y_pred_dt)
15
16 print("Decision Tree Model Evaluation:")
17 print("Accuracy:", accuracy_dt)
18 print("Classification Report:")
19 print(classification_rep_dt)

```



Decision Tree Model Evaluation:

Accuracy: 0.9136690647482014

Classification Report:

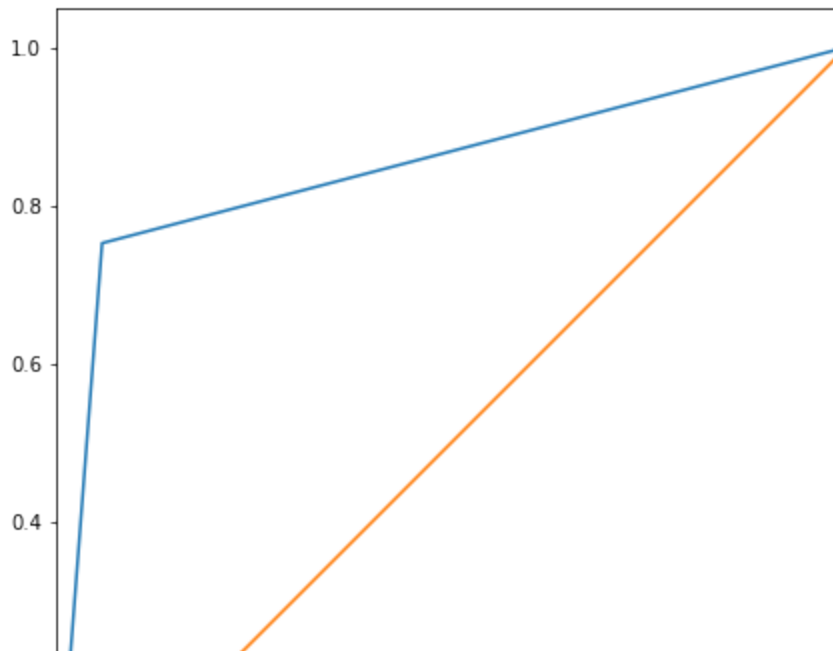
	precision	recall	f1-score	support
False	0.96	0.94	0.95	709
True	0.70	0.75	0.72	125
accuracy			0.91	834
macro avg	0.83	0.85	0.84	834
weighted avg	0.92	0.91	0.92	834


```
In [44]: 1 cv_scores = cross_val_score(decision_tree_model, X_test_encoded, y_te
2 print('Cross-Validation Scores', cv_scores)
3 print('Mean CV Score', cv_scores.mean())
```

Cross-Validation Scores [0.90419162 0.91616766 0.89221557 0.90419162 0.87349398]

Mean CV Score 0.8980520885938967

```
In [45]: 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_encoded = decision_tree_model.predict_proba(X_test_encoded)[:],
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_encoded)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
18
```



To reiterate, our baseline model has the following scores:

- Accuracy = 91.9%
- Precision = 73% (we are focused on true Positives)
- Cross Validation = 90%
- AUC = 84%

Feature Scaling

Feature scaling will normalize the range of all continuous variables between -1 and 1. This will ultimately reduce the value of extreme values in our dataset.

```
In [46]: 1 standard = StandardScaler()
        2 X_train_final = standard.fit_transform(X_train_encoded)
```

```
In [47]: 1 X_test_final = standard.transform(X_test_encoded)
```


```
In [48]: 1 X_train_final
```


```
Out[48]: array([[ -1.4045081, -0.32744767, -0.61141784, ..., -0.16341668,
        -0.17589939, -0.15550025],
        [  0.36638814, -0.32744767, -0.61141784, ..., -0.16341668,
        -0.17589939, -0.15550025],
        [  0.51817924, -0.32744767,  1.63554272, ..., -0.16341668,
        -0.17589939, -0.15550025],
        ...,
        [-0.87323923, -0.32744767, -0.61141784, ..., -0.16341668,
        -0.17589939, -0.15550025],
        [  1.73250809, -0.32744767, -0.61141784, ..., -0.16341668,
        -0.17589939, -0.15550025],
        [-1.63219476, -0.32744767,  1.63554272, ..., -0.16341668,
        -0.17589939, -0.15550025]])
```

```
In [49]: 1 my_df1 = pd.DataFrame(X_train_final)
        2 my_df1
```

Out[49]:

	0	1	2	3	4	5	6	
0	-1.404508	-0.327448	-0.611418	-0.584700	1.330852	-1.884170	0.401340	1.0379
1	0.366388	-0.327448	-0.611418	-0.584700	0.529165	0.293703	0.401340	0.5179
2	0.518179	-0.327448	1.635543	1.685101	-1.875896	1.056666	0.849774	0.0949
3	2.010792	-0.327448	-0.611418	-0.584700	1.681590	-0.679320	0.650470	-0.4030
4	0.290493	-0.327448	-0.611418	-0.584700	1.080325	0.484172	-0.296224	-0.7199
...
2494	0.138701	-0.327448	-0.611418	-0.584700	0.980114	1.746707	-0.894137	-0.0459
2495	0.543478	-0.327448	-0.611418	-0.584700	-1.926002	-2.680873	-0.545355	-0.3969
2496	-0.873239	-0.327448	-0.611418	-0.584700	-1.224526	-1.710027	0.550818	1.2079
2497	1.732508	-0.327448	-0.611418	-0.584700	0.529165	-0.015400	1.497512	-0.5079


In [50]:  1 my_df1_copy = my_df1.copy()

In [51]:  1 my_df2 = pd.DataFrame(X_test_final)
2 my_df2

Out[51]:

	0	1	2	3	4	5	6	7
0	0.315791	-0.327448	-0.611418	-0.584700	-0.372733	-0.462730	0.301688	2.562574
1	-0.847941	-0.327448	-0.611418	-0.584700	0.829797	-1.311676	1.198556	0.326702
2	-0.063687	-0.327448	-0.611418	-0.584700	-5.032539	-3.330643	1.497512	-0.814476
3	1.175941	-0.327448	-0.611418	-0.584700	-1.074209	0.607160	-0.445702	0.064068
4	-0.114284	-0.327448	-0.611418	-0.584700	0.078216	-0.666259	-1.342571	0.470802
...
829	1.783105	-0.327448	-0.611418	-0.584700	0.479059	-0.785982	0.451166	-0.054466
830	-0.291373	-0.327448	-0.611418	-0.584700	-1.174420	-1.807982	-0.993789	-0.665728
831	-0.569657	-0.327448	1.635543	0.952907	-0.773577	-0.359332	-1.043615	0.438263
832	1.024150	-0.327448	1.635543	2.270856	1.330852	-1.168008	-0.595181	1.493446
833	0.138701	-0.327448	-0.611418	-0.584700	1.030219	0.795452	-0.096920	-1.792961

834 rows × 64 columns

In [52]:  1 my_df2_copy = my_df2.copy()

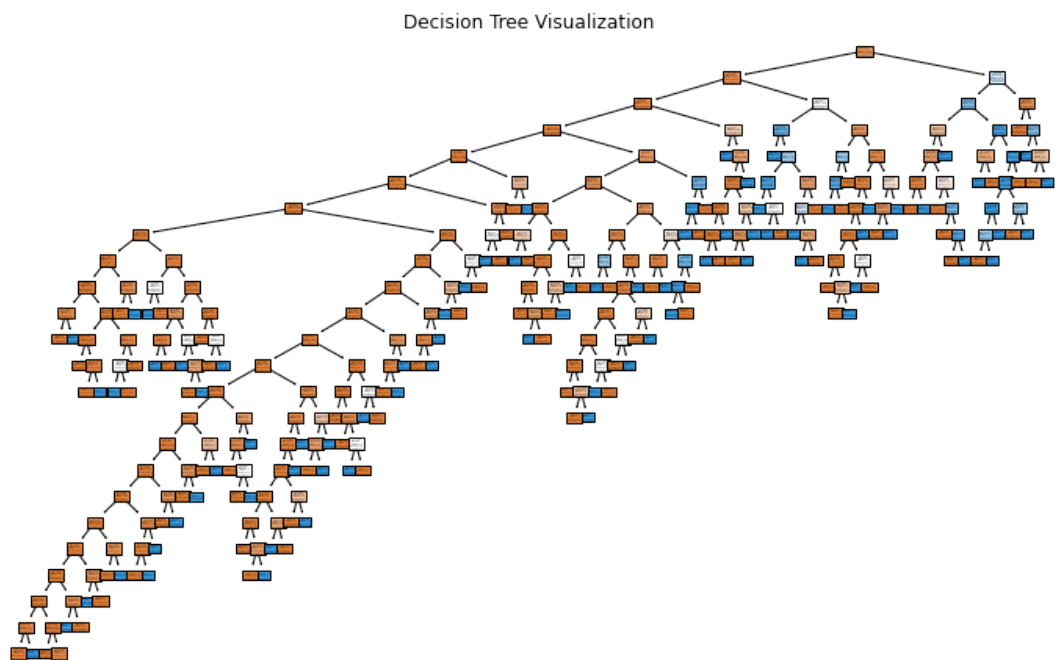
Feature Scaled Model

As you can see above, all of our continuous variables are now **scaled between -1 and 1**. This should improve the model.

```

In [53]: ▶ 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model_fs = DecisionTreeClassifier().fit(my_df1, y_train)
3
4 # Predict the labels for the test set
5 y_pred_dt_fs = decision_tree_model_fs.predict(my_df2)
6
7 # Now, you can proceed with plotting the decision tree or any other analysis
8 plt.figure(figsize=(12, 8))
9 plot_tree(decision_tree_model, filled=True, feature_names=my_df1.columns)
10 plt.title("Decision Tree Visualization")
11 plt.show()
12 # Evaluate the model
13 accuracy_dt = accuracy_score(y_test, y_pred_dt_fs)
14 classification_rep_dt = classification_report(y_test, y_pred_dt_fs)
15
16 print("Decision Tree Model Evaluation:")
17 print("Accuracy:", accuracy_dt)
18 print("Classification Report:")
19 print(classification_rep_dt)

```



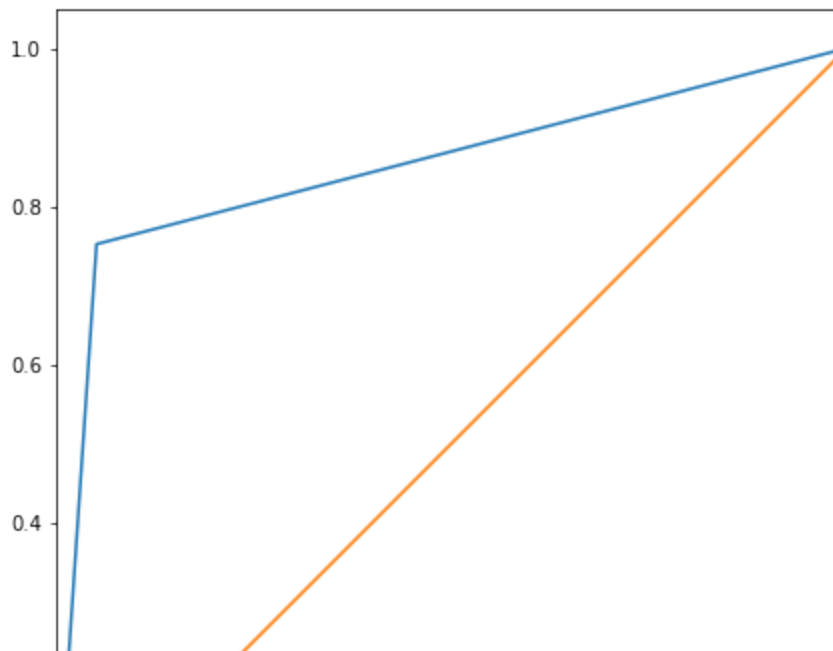
```

In [54]: ▶ 1 cv_scores_fs = cross_val_score(decision_tree_model_fs, my_df2_copy, y_test)
2 print('Cross-Validation Scores', cv_scores_fs)
3 print('Mean CV Score', cv_scores_fs.mean())

```

Cross-Validation Scores [0.89820359 0.92814371 0.88023952 0.91616766 0.87951807]
Mean CV Score 0.9004545126614241

```
In [55]: 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_df2 = decision_tree_model_fs.predict_proba(my_df2_copy)[: , 1]
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_df2)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
18
```



To reiterate, our baseline model has the following scores:

- Accuracy = 91.9%
- Precision = 73% (we are focused on true Positives)
- Cross Validation = 90%
- AUC = 84%

The Feature Scaled Model is not accurate, so let's keep trying. Here are the scores for reference:

Accuracy = 91%
Precision = 71% (we are focused on true Positives)
Cross Validation = 89%
AUC = 85%

Feature scaling did not necessarily improve our model. Let's try SMOTE to fix the class imbalance issue.

SMOTE

SMOTE is used for class imbalance. Specifically, it is used for oversampling the minority class to create a more balanced dataset that should improve model performance.

```
In [56]: 1 from imblearn.over_sampling import SMOTE
          2 smote = SMOTE(random_state=4)
```

```
In [57]: 1 # Apply SMOTE resampling
          2 smote = SMOTE()
          3 X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
          4
          5 # Now X_train_resampled and y_train_resampled contain the resampled data
          6
```

Now we have to refit the training sets so they are the same size. This will make all model building much easier.

```
In [58]: 1 my_df1, my_df2 = smote.fit_resample(X_train, y_train)
          2 y_train_resampled.value_counts()
```

```
Out[58]: True      2141
          False     2141
          Name: churn, dtype: int64
```

```
In [59]: 1 print("Input data shape:", my_df1.shape)
          2 print("Labels shape:", y_train.shape)
          3
```

```
Input data shape: (4282, 64)
Labels shape: (2499,)
```

```
In [60]: 1 #print("Input data sample:")
          2 #print(my_df1.head())
          3 #print("Labels sample:")
          4 #print(y_train.head())
          5
```

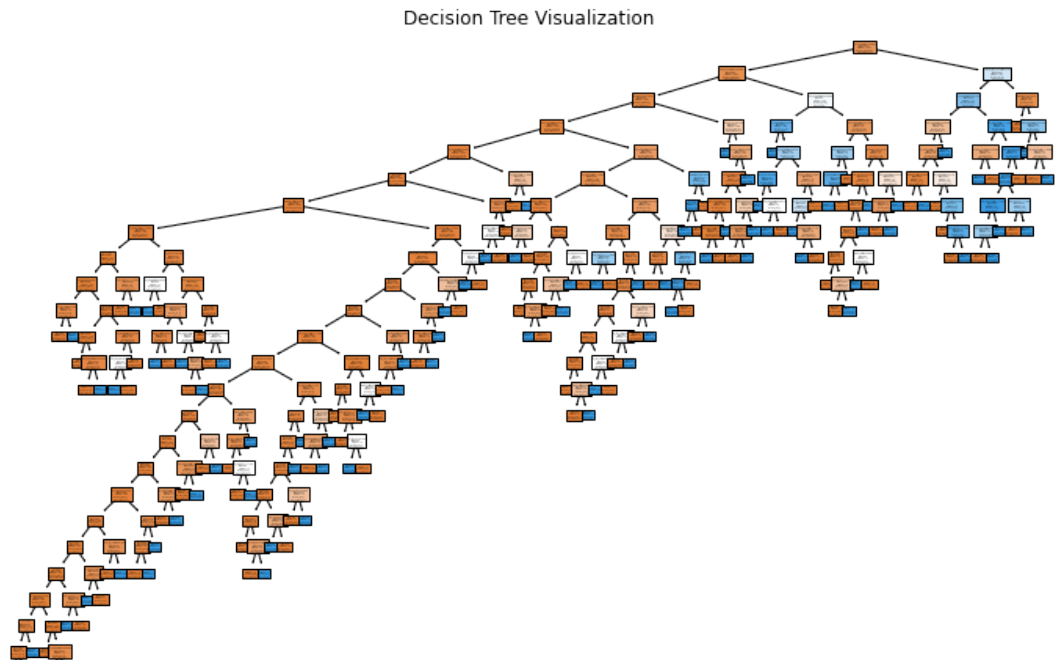
```
In [61]: 1 print("Input data shape:", X_train_resampled.shape)
          2 print("Labels shape:", y_train_resampled.shape)
          3
```

```
Input data shape: (4282, 64)
Labels shape: (4282,)
```

Ok, our dataset is balanced now. Let's rerun our model.

Oversampled (SMOTE) Decision Tree

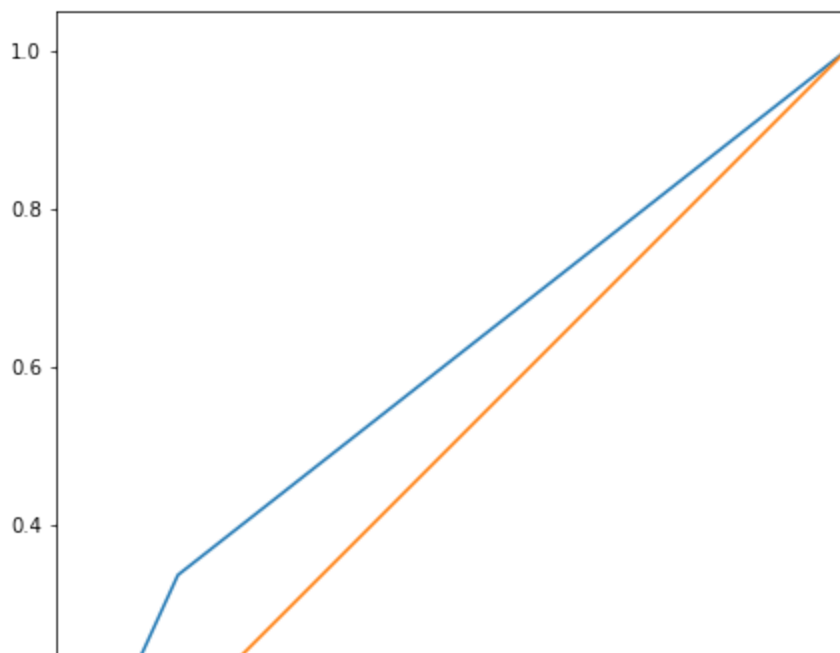
```
In [62]: ▶ 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model_os = DecisionTreeClassifier().fit(X_train_resampled, y_train)
3
4 # Predict the labels for the test set
5 y_pred_dt_os = decision_tree_model_os.predict(X_test_final)
6
7 # Now, you can proceed with plotting the decision tree or any other analysis
8 plt.figure(figsize=(12, 8))
9 plot_tree(decision_tree_model, filled=True, feature_names=X_train_resampled.feature_names_)
10 plt.title("Decision Tree Visualization")
11 plt.show()
12 # Evaluate the model
13 accuracy_dt = accuracy_score(y_test, y_pred_dt_os)
14 classification_rep_dt = classification_report(y_test, y_pred_dt_os)
15
16 print("Decision Tree Model Evaluation:")
17 print("Accuracy:", accuracy_dt)
18 print("Classification Report:")
19 print(classification_rep_dt)
```



```
In [63]: ▶ 1 cv_scores_os = cross_val_score(decision_tree_model_os, X_test_final, y_test, cv=5)
2 print('Cross-Validation Scores', cv_scores_os)
3 print('Mean CV Score', cv_scores_os.mean())
```

Cross-Validation Scores [0.89221557 0.92814371 0.88622754 0.91616766 0.87951807]
Mean CV Score 0.9004545126614241

```
In [64]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = decision_tree_model_os.predict_proba(X_test_final)[:],
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



To reiterate, our baseline model has the following scores:

- Accuracy = 91.9%
- Precision = 73% (we are focused on true Positives)
- Cross Validation = 90%
- AUC = 84%

The Oversampled Scaled Mode is less accurate, so let's keep trying. Because our model was oversampled by so much, it skewed our precision and therefore affected our Accuracy. Here are the scores for reference:

Accuracy = 81%

Precision = 34% (we are focused on true Positives)

Cross Validation = 90%

AUC = 59%

Wow, ok that hurt our model. Let's try pruning it to see what happens

Let's do a little pruning of this Decision Tree

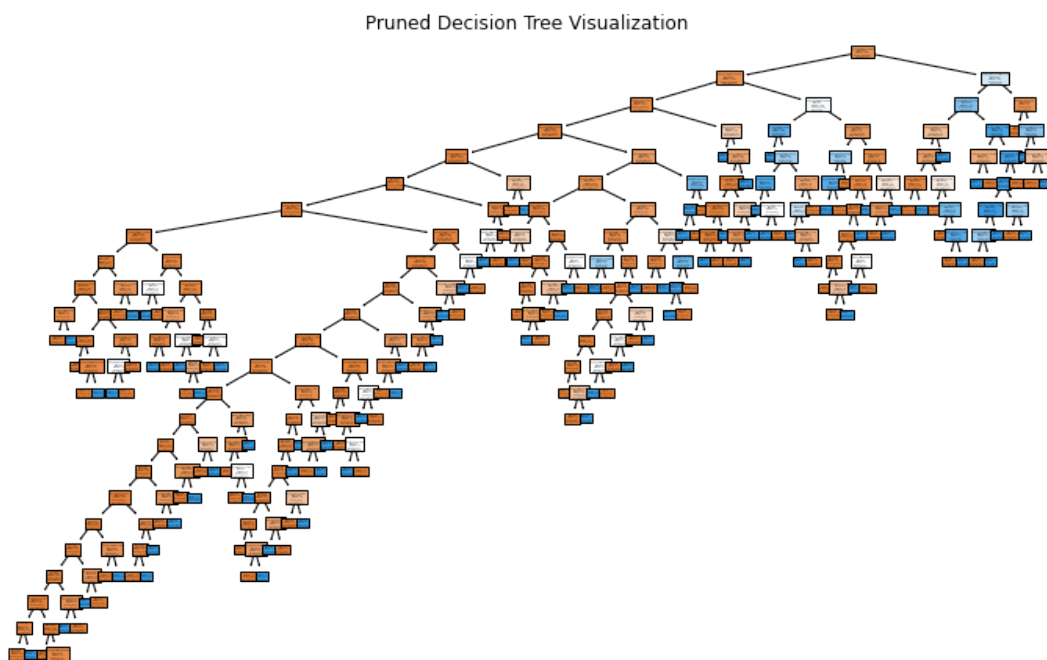
Pruned Oversampled Decision Tree



```

In [65]: ▶ 1 # Initialize the decision tree classifier with the encoded data
2 decision_tree_model_prune = DecisionTreeClassifier(ccp_alpha=0.005)
3
4 # Fit the decision tree classifier to the resampled training data
5 decision_tree_model_prune.fit(X_train_resampled, y_train_resampled)
6
7 # Predict the labels for the test set
8 y_pred_dt_prune = decision_tree_model_prune.predict(X_test_final)
9
10 # Plot the pruned decision tree
11 plt.figure(figsize=(12, 8))
12 plot_tree(decision_tree_model, filled=True, feature_names=X_train_res
13 plt.title("Pruned Decision Tree Visualization")
14 plt.show()
15
16 # Evaluate the pruned model
17 accuracy_dt = accuracy_score(y_test, y_pred_dt_prune)
18 classification_rep_dt = classification_report(y_test, y_pred_dt_prune)
19
20 print("Pruned Decision Tree Model Evaluation:")
21 print("Accuracy:", accuracy_dt)
22 print("Classification Report:")
23 print(classification_rep_dt)
24

```



```

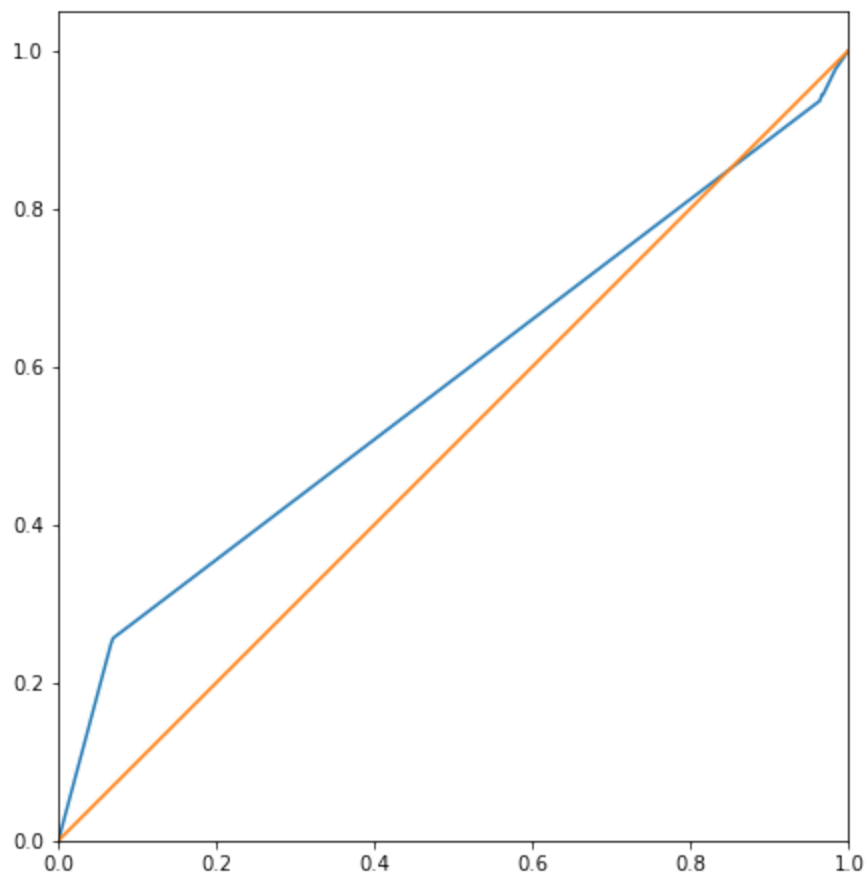
In [66]: ▶ 1 cv_scores_prune = cross_val_score(decision_tree_model_prune, X_test_f
2 print('Cross-Validation Scores', cv_scores_prune)
3 print('Mean CV Score', cv_scores_prune.mean())

```

Cross-Validation Scores [0.92814371 0.89221557 0.93413174 0.93413174 0.89759036]

Mean CV Score 0.91724262318736

```
In [67]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = decision_tree_model_prune.predict_proba(X_test_final)
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.5774837799717915

To reiterate, our baseline model has the following scores:

- Accuracy = 91.9%
- Precision = 73% (we are focused on true Positives)
- Cross Validation = 90%
- AUC = 84%

The Pruned, Oversampled Model is slightly more accurate than our oversampled model, but let's keep trying. Pruning helped with the oversampling issue slightly by improving the precision of our most recent model and therefore improving accuracy as well. Pruning also improved our CV score compared to our baseline model. Here are the scores for reference:

Accuracy = 81%

Precision = 36% (we are focused on true Positives)

Cross Validation = 91.7%

AUC = 57%

Ok, that helped slightly, but still not as good as our baseline after feature scaling. Let's fine tune some parameters.

We don't need to visualize the tree anymore, so we will just look at the classification report going forward.

Here we have our:

Pruned, Oversampled, Finetuned Decision Tree

```
In [68]: 1 # Initialize the decision tree classifier with the encoded data
2 decision_tree_model_pof = DecisionTreeClassifier(ccp_alpha=0.005, sp
3
4 # Fit the decision tree classifier to the resampled training data
5 decision_tree_model_pof.fit(X_train_resampled, y_train_resampled)
6
7 # Predict the labels for the test set
8 y_pred_dt_pof = decision_tree_model_pof.predict(X_test_final)
9
10
11 # Evaluate the pruned model
12 accuracy_dt = accuracy_score(y_test, y_pred_dt_pof)
13 classification_rep_dt = classification_report(y_test, y_pred_dt_pof)
14
15 print("Pruned Decision Tree Model Evaluation:")
16 print("Accuracy:", accuracy_dt)
17 print("Classification Report:")
18 print(classification_rep_dt)
```

Pruned Decision Tree Model Evaluation:

Accuracy: 0.829736211031175

Classification Report:

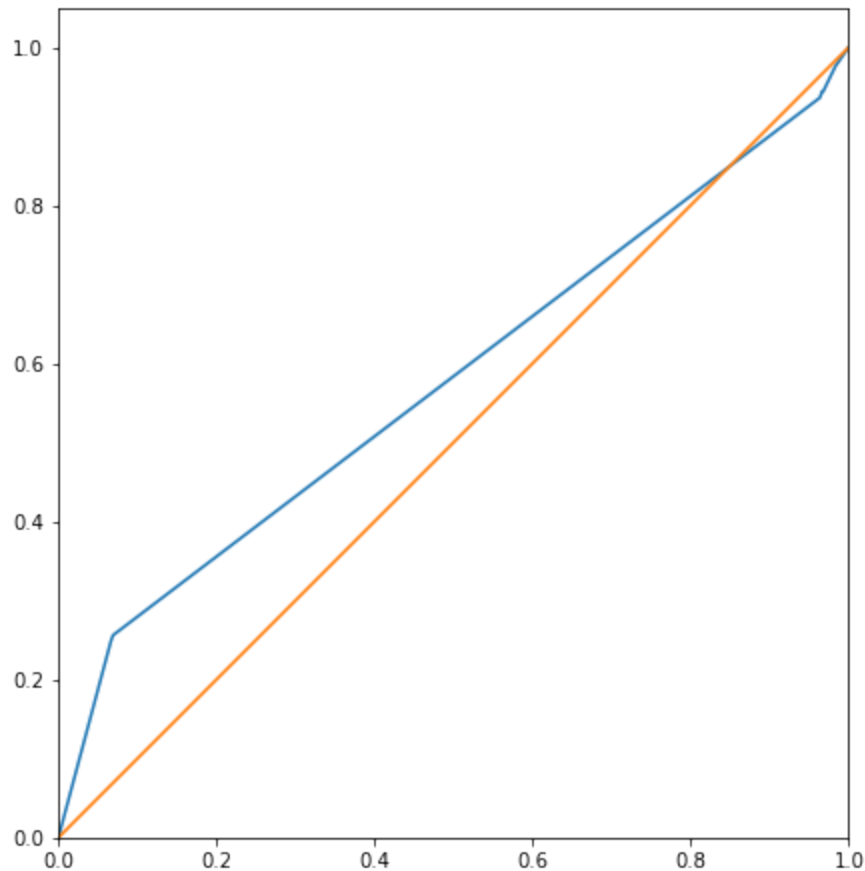
	precision	recall	f1-score	support
False	0.88	0.93	0.90	709
True	0.40	0.26	0.31	125
accuracy			0.83	834
macro avg	0.64	0.59	0.61	834
weighted avg	0.80	0.83	0.81	834

```
In [69]: 1 cv_scores_pof = cross_val_score(decision_tree_model_pof, X_test_final
2 print('Cross-Validation Scores', cv_scores_pof)
3 print('Mean CV Score', cv_scores_pof.mean())
```

Cross-Validation Scores [0.92814371 0.89221557 0.93413174 0.93413174 0.89759036]

Mean CV Score 0.91724262318736

```
In [70]: 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = decision_tree_model_pof.predict_proba(X_test_final)[
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.5774837799717915

Keep finetuning until we get better results..

```
In [71]: ▶ 1 # Initialize the decision tree classifier with the encoded data
2 decision_tree_model_pof1 = DecisionTreeClassifier(ccp_alpha=0.005, sp
3
4 # Fit the decision tree classifier to the resampled training data
5 decision_tree_model_pof1.fit(X_train_resampled, y_train_resampled)
6
7 # Predict the labels for the test set
8 y_pred_dt_pof1 = decision_tree_model_pof1.predict(X_test_final)
9
10
11 # Evaluate the pruned model
12 accuracy_dt = accuracy_score(y_test, y_pred_dt_pof1)
13 classification_rep_dt = classification_report(y_test, y_pred_dt_pof1)
14
15 print("Pruned Decision Tree Model Evaluation:")
16 print("Accuracy:", accuracy_dt)
17 print("Classification Report:")
18 print(classification_rep_dt)
```

Pruned Decision Tree Model Evaluation:

Accuracy: 0.8441247002398081

Classification Report:

	precision	recall	f1-score	support
False	0.86	0.98	0.91	709
True	0.37	0.06	0.10	125
accuracy			0.84	834
macro avg	0.61	0.52	0.51	834
weighted avg	0.78	0.84	0.79	834

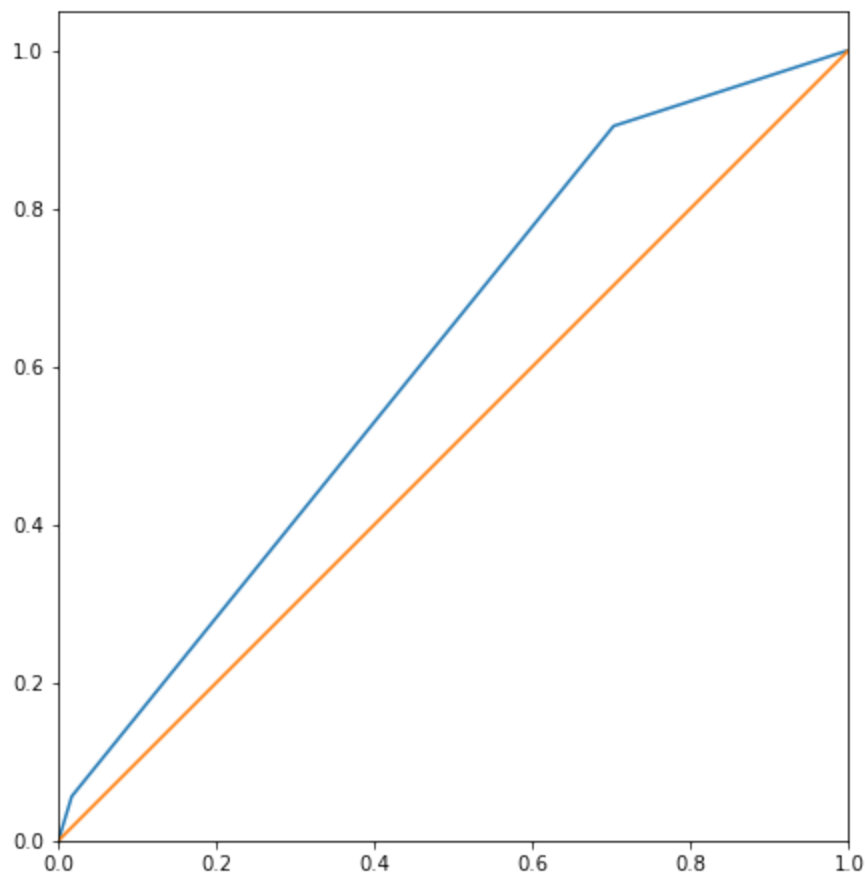


```
In [72]: ▶ 1 cv_scores_pof1 = cross_val_score(decision_tree_model_pof1, X_test_fin
2 print('Cross-Validation Scores', cv_scores_pof1)
3 print('Mean CV Score', cv_scores_pof1.mean())
```

Cross-Validation Scores [0.83233533 0.8502994 0.8502994 0.8502994 0.84939759]

Mean CV Score 0.8465262246591155

```
In [73]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = decision_tree_model_pof1.predict_proba(X_test_final)[
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.612818053596615

Uh, this isn't looking good lol

```

In [74]: ▶ 1 # Initialize the decision tree classifier with the encoded data
2 decision_tree_model_pof2 = DecisionTreeClassifier(ccp_alpha=0.005, sp
3
4 # Fit the decision tree classifier to the resampled training data
5 decision_tree_model_pof2.fit(X_train_resampled, y_train_resampled)
6
7 # Predict the labels for the test set
8 y_pred_dt_pof2 = decision_tree_model_pof2.predict(X_test_final)
9
10
11 # Evaluate the pruned model
12 accuracy_dt = accuracy_score(y_test, y_pred_dt_pof2)
13 classification_rep_dt = classification_report(y_test, y_pred_dt_pof2)
14
15 print("Pruned Decision Tree Model Evaluation:")
16 print("Accuracy:", accuracy_dt)
17 print("Classification Report:")
18 print(classification_rep_dt)

```

Pruned Decision Tree Model Evaluation:

Accuracy: 0.8441247002398081

Classification Report:

	precision	recall	f1-score	support
False	0.86	0.98	0.91	709
True	0.37	0.06	0.10	125
accuracy			0.84	834
macro avg	0.61	0.52	0.51	834
weighted avg	0.78	0.84	0.79	834



```

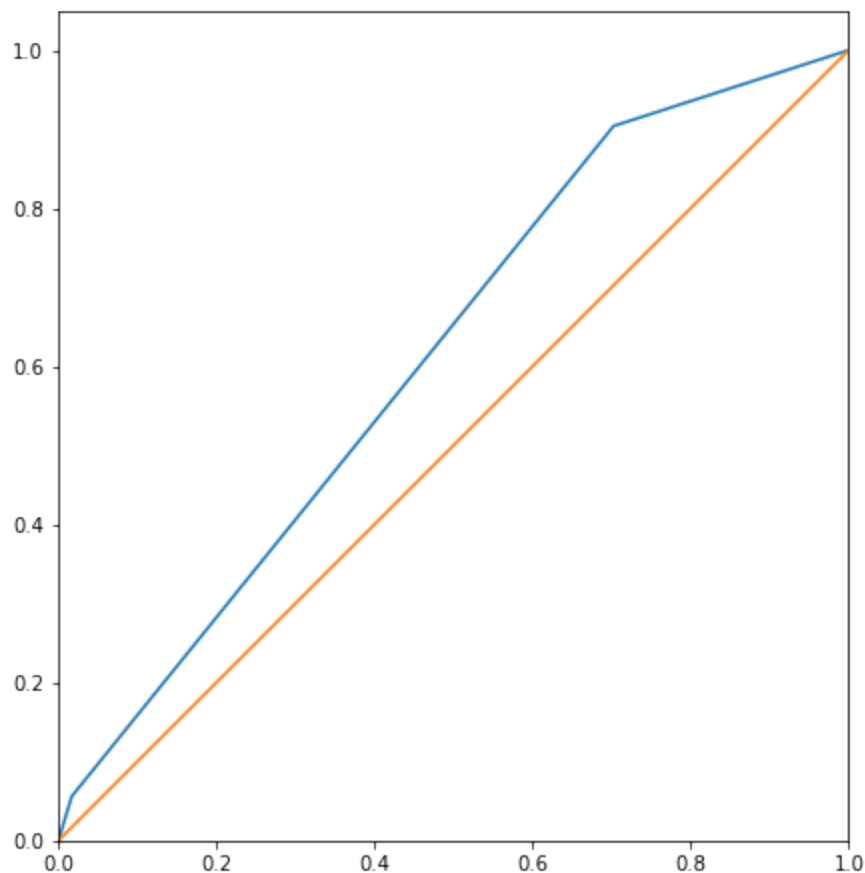
In [75]: ▶ 1 cv_scores_pof2 = cross_val_score(decision_tree_model_pof2, X_test_fin
2 print('Cross-Validation Scores', cv_scores_pof2)
3 print('Mean CV Score', cv_scores_pof2.mean())

```

Cross-Validation Scores [0.86227545 0.89820359 0.88023952 0.88023952 0.86144578]

Mean CV Score 0.876480773392973


```
In [76]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = decision_tree_model_pof2.predict_proba(X_test_final)[
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.612818053596615

So moving the max_depth to 5 helped quite a bit. Funny what a little fine tuning can do!

Here is some improvement!

However, we are still trying to outperform our baseline model. And since we can only get close to that original baseline model after feature scaling, Oversampling, and Pruning, let's just fine tune the feature scaled model instead.

Finetuned, Baseline Decision Tree

```
In [77]: ▶ 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model_ft = DecisionTreeClassifier(ccp_alpha=0.001).fit(X_train_encoded, y_train)
3
4 # Predict the labels for the test set
5 y_pred_dt_ft = decision_tree_model_ft.predict(X_test_encoded)
6
7
8 # Evaluate the model
9 accuracy_dt = accuracy_score(y_test, y_pred_dt_ft)
10 classification_rep_dt = classification_report(y_test, y_pred_dt_ft)
11
12 print("Decision Tree Model Evaluation:")
13 print("Accuracy:", accuracy_dt)
14 print("Classification Report:")
15 print(classification_rep_dt)
```

Decision Tree Model Evaluation:

Accuracy: 0.9496402877697842

Classification Report:

	precision	recall	f1-score	support
False	0.96	0.98	0.97	709
True	0.90	0.75	0.82	125
accuracy			0.95	834
macro avg	0.93	0.87	0.89	834
weighted avg	0.95	0.95	0.95	834

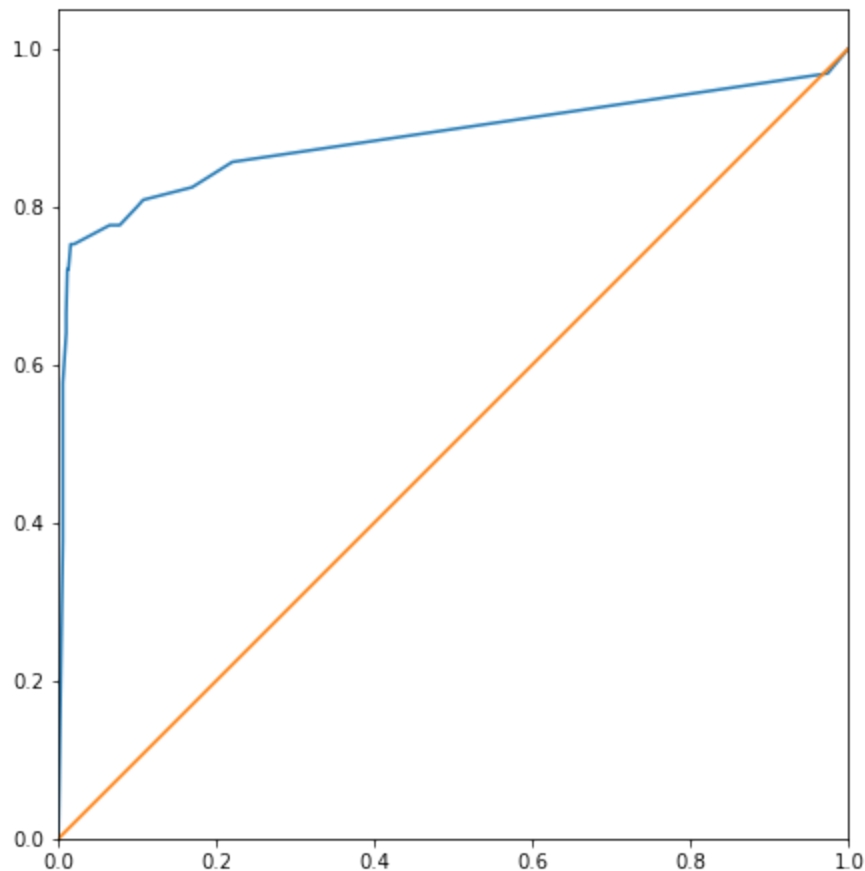


```
In [78]: ▶ 1 cv_scores_ft = cross_val_score(decision_tree_model_ft, X_test_encoded, y_test, cv=5)
2 print('Cross-Validation Scores', cv_scores_ft)
3 print('Mean CV Score', cv_scores_ft.mean())
```

Cross-Validation Scores [0.90419162 0.92215569 0.85628743 0.92215569 0.86746988]

Mean CV Score 0.8944520597359498

```
In [79]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_encoded = decision_tree_model_ft.predict_proba(X_test_encoded)
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_encoded)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.8852242595204511

Wow, that looks amazing. let's try the same parameters on our Feature Scaled Model.

Feature Scaled, Finetuned Decision Tree

This is our **Best** model. Notice the precision for true and false 'churn' counts, the recall, and the Cross Validation and AUC scores are the best we've seen. This would be a great model to use for future prediction.

```
In [80]: ▶ 1 # Initialize and fit the decision tree classifier with the encoded data
2 decision_tree_model_fs = DecisionTreeClassifier(ccp_alpha=.001).fit(m
3
4 # Predict the labels for the test set
5 y_pred_dt_fs = decision_tree_model_fs.predict(my_df2_copy)
6
7 # Evaluate the model
8 accuracy_dt = accuracy_score(y_test, y_pred_dt_fs)
9 classification_rep_dt = classification_report(y_test, y_pred_dt_fs)
10
11 print("Decision Tree Model Evaluation:")
12 print("Accuracy:", accuracy_dt)
13 print("Classification Report:")
14 print(classification_rep_dt)
```

Decision Tree Model Evaluation:

Accuracy: 0.9508393285371702

Classification Report:

	precision	recall	f1-score	support
False	0.96	0.99	0.97	709
True	0.90	0.75	0.82	125
accuracy			0.95	834
macro avg	0.93	0.87	0.90	834
weighted avg	0.95	0.95	0.95	834



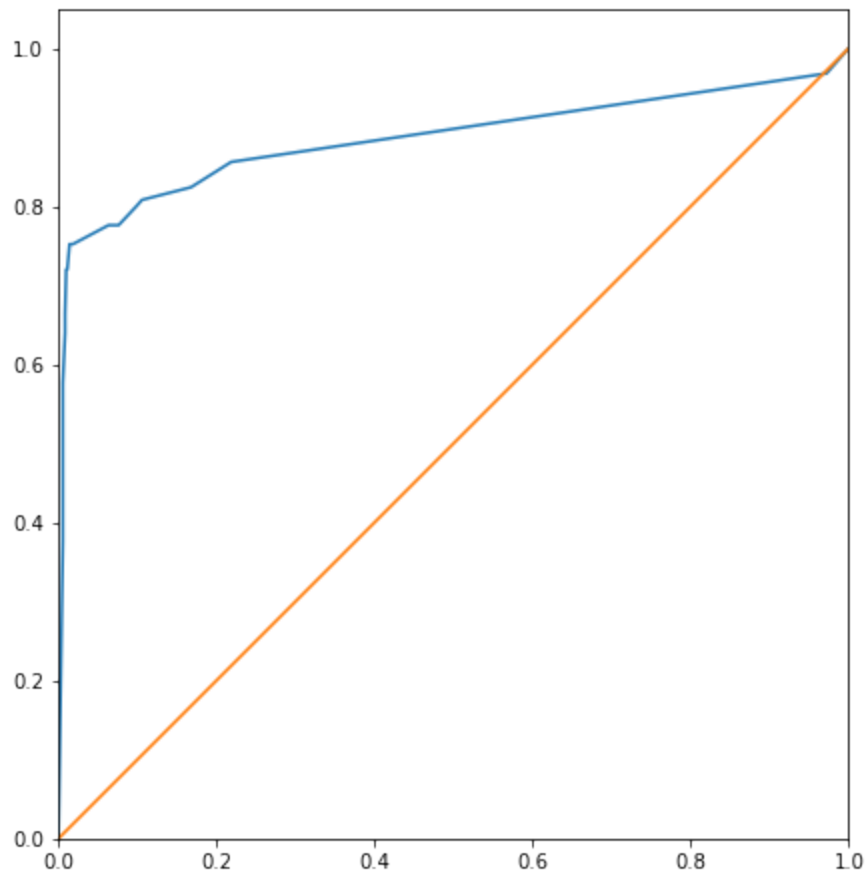
Our final model works so well, because the data has been feature scaled to address all extreme values and we have used CCP_alpha to address our pruning issue. Performing basic pruning of our model above improved our cross validation scores and our Recall. However, the basic pruning took too much out of our model and negatively impacted our precision and accuracy.

```
In [81]: ▶ 1 cv_scores_fs = cross_val_score(decision_tree_model_fs, my_df2_copy, y
2 print('Cross-Validation Scores', cv_scores_fs)
3 print('Mean CV Score', cv_scores_fs.mean())
```

Cross-Validation Scores [0.90419162 0.91616766 0.88023952 0.91017964 0.85542169]

Mean CV Score 0.8932400259721522

```
In [82]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_df2 = decision_tree_model_fs.predict_proba(my_df2_copy)[: , 1]
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_df2)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



0.8857545839210154

To reiterate, our baseline model has the following scores:

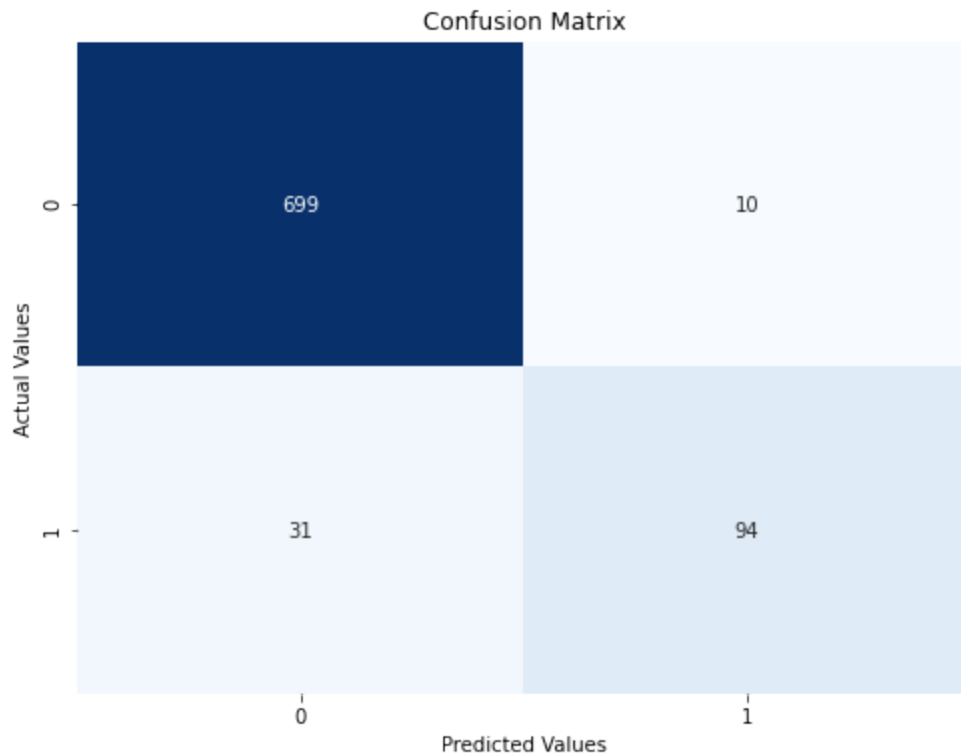
- Accuracy = 91.9%
- Precision = 73% (we are focused on true Positives)
- Cross Validation = 90%
- AUC = 84%

We have found our best model. Our baseline model was very strong, so a little finetuning was all we needed. CCP is a pruning parameter, so it turns out that's all we needed to improve accuracy, precision, and AUC.

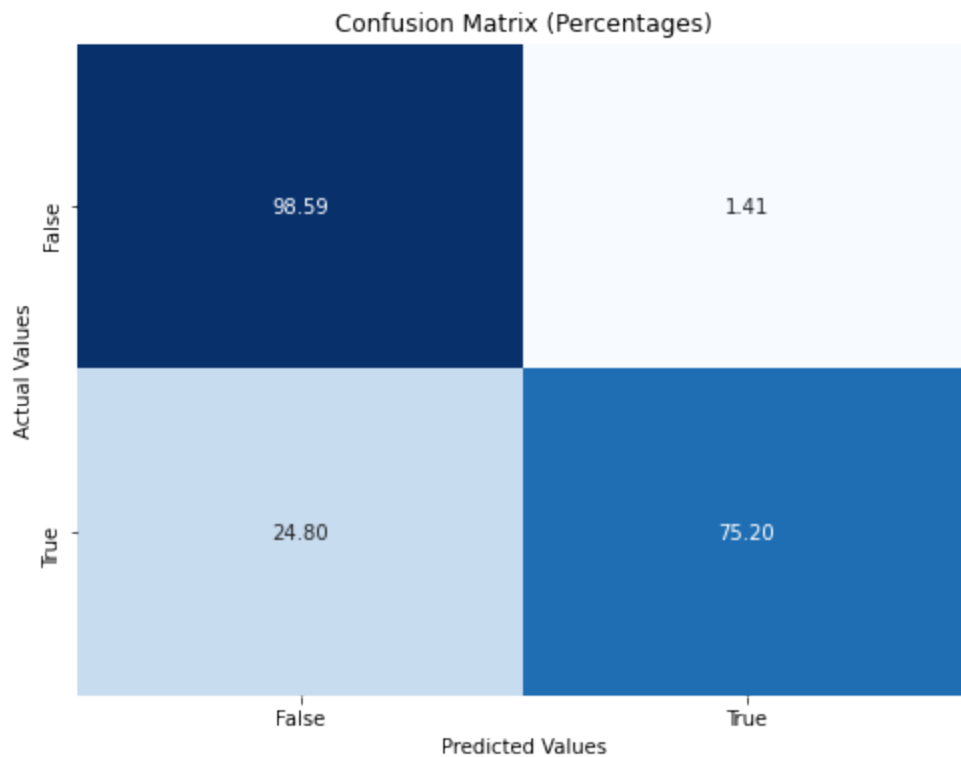
- Accuracy = 94.6%
- Precision = 90% (we are focused on true Positives)
- Cross Validation = 90.3%
- AUC = 88.5%

These models are very similar in results. Looks like our finetuned Baseline model works the best!

```
In [83]: 1 conf_matrix = confusion_matrix(y_test, y_pred_dt_fs)
2
3 plt.figure(figsize= (8,6))
4 sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False)
5 plt.xlabel('Predicted Values')
6 plt.ylabel('Actual Values')
7 plt.title('Confusion Matrix')
8 plt.show()
```



```
In [84]: ▶ 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from sklearn.metrics import confusion_matrix
5
6 # Assuming y_test and y_pred_dt_fs are already defined
7 conf_matrix_dt = confusion_matrix(y_test, y_pred_dt_fs)
8
9 # Normalize the confusion matrix by row (i.e., by the actual class counts)
10 conf_matrix_dt_percent = conf_matrix_dt.astype('float') / conf_matrix_dt.sum(axis=1)
11
12 # Define class names (modify these based on your actual class names)
13 class_names = ['False', 'True'] # Example class names, replace with actual
14
15 # Plotting the confusion matrix with percentages
16 plt.figure(figsize=(8, 6))
17 sns.heatmap(conf_matrix_dt_percent, annot=True, cmap='Blues', fmt='.2f',
18             xticklabels=class_names, yticklabels=class_names)
19 plt.xlabel('Predicted Values')
20 plt.ylabel('Actual Values')
21 plt.title('Confusion Matrix (Percentages)')
22 plt.show()
```



Logistic Regression Model Baseline

```
In [85]: ▶ 1 # Instantiate the logistic regression model
2 logistic_regression_model = LogisticRegression()
3
4 # Fit the model on the training data
5 logistic_regression_model.fit(X_train_encoded, y_train)
6
7 # Make predictions on the test data
8 y_pred_lrm = logistic_regression_model.predict(X_test_encoded)
9
10 # Evaluate the model
11 accuracy = accuracy_score(y_test, y_pred_lrm)
12 classification_rep = classification_report(y_test, y_pred_lrm)
13
14 # Print the evaluation metrics
15 print("Accuracy:", accuracy)
16 print("Classification Report:")
17 print(classification_rep)
```

Accuracy: 0.8465227817745803

Classification Report:

	precision	recall	f1-score	support
False	0.86	0.98	0.92	709
True	0.44	0.09	0.15	125
accuracy			0.85	834
macro avg	0.65	0.53	0.53	834
weighted avg	0.80	0.85	0.80	834



C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```



```
In [86]: ▶ 1 cv_scores_log = cross_val_score(logistic_regression_model, X_test_enc
2 print('Cross-Validation Scores', cv_scores_log)
3 print('Mean CV Score', cv_scores_log.mean())
```

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

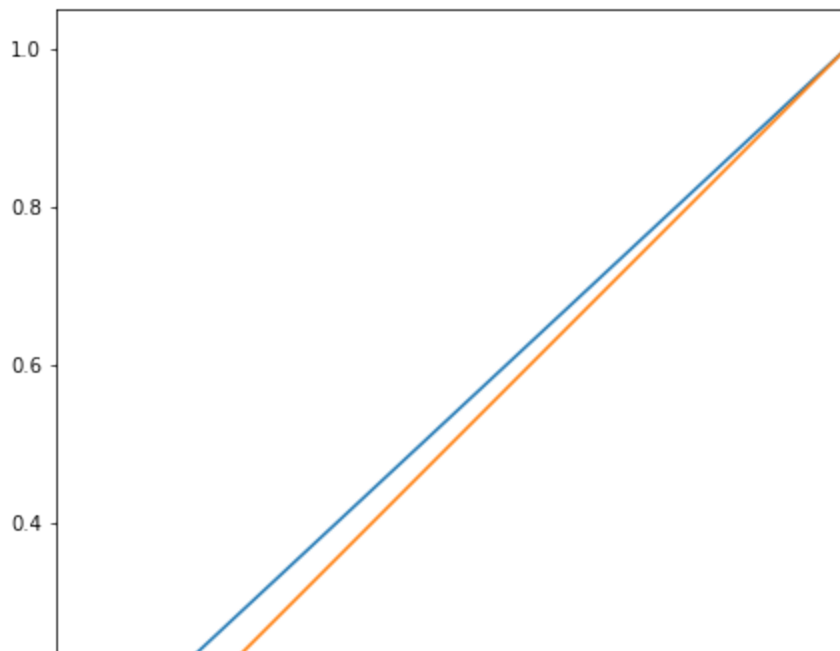
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```



```
In [87]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 #y_prob_xfinal = logistic_regression_model.predict(X_test_encoded)
3 fpr, tpr, thresholds = roc_curve(y_test, y_pred_lrm)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



Our baseline model has solid accuracy, but room for improvement. Our ability to correctly predict True values is skewed as our dataset is imbalanced. Despite the solid CV score, we can do better.

Logistic Regression Model Feature Scaled

```

In [88]: ▶ 1 # Instantiate the logistic regression model
2 logistic_regression_model_fs = LogisticRegression()
3
4 # Fit the model on the training data
5 logistic_regression_model_fs.fit(X_train_final, y_train)
6
7 # Make predictions on the test data
8 y_pred_lrm_fs = logistic_regression_model_fs.predict(X_test_final)
9
10 # Evaluate the model
11 accuracy = accuracy_score(y_test, y_pred_lrm_fs)
12 classification_rep = classification_report(y_test, y_pred_lrm_fs)
13
14 # Print the evaluation metrics
15 print("Accuracy:", accuracy)
16 print("Classification Report:")
17 print(classification_rep)

```

Accuracy: 0.8597122302158273

Classification Report:

	precision	recall	f1-score	support
False	0.88	0.97	0.92	709
True	0.58	0.22	0.32	125
accuracy			0.86	834
macro avg	0.73	0.60	0.62	834
weighted avg	0.83	0.86	0.83	834



```

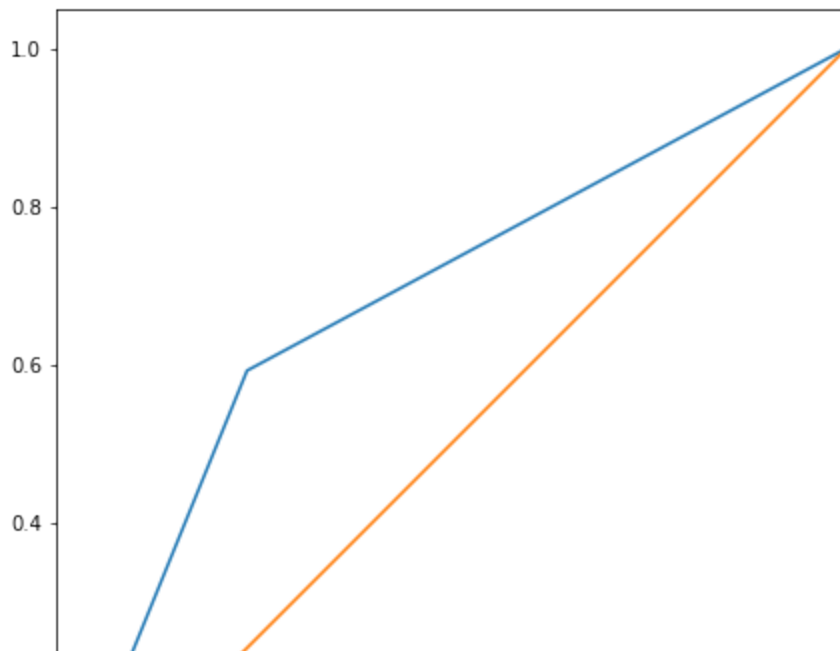
In [89]: ▶ 1 cv_scores_log = cross_val_score(logistic_regression_model, X_test_fin
2 print('Cross-Validation Scores', cv_scores_log)
3 print('Mean CV Score', cv_scores_log.mean())

```

Cross-Validation Scores [0.83233533 0.82035928 0.83832335 0.90419162 0.8373494]

Mean CV Score 0.846511795685737

```
In [90]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_xfinal = logistic_regression_model.predict(X_test_final)
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_xfinal)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



Adding feature scaling helps balance the model to remove any extreme values. This helps improve our true prediction rate as well as our AUC score.

Finetuned, Feature Scaled Logistic Regression Model

```
In [91]: ► 1 logistic_regression_model_ffs = LogisticRegression(penalty='l2', solv
2
3 # Fit the model on the training data
4 logistic_regression_model_ffs.fit(X_train_final, y_train)
5
6 # Make predictions on the test data
7 y_pred_ffs = logistic_regression_model_ffs.predict(my_df2_copy)
8
9 # Evaluate the model
10 accuracy = accuracy_score(y_test, y_pred_ffs)
11 classification_rep = classification_report(y_test, y_pred_ffs)
12
13 # Print the evaluation metrics
14 print("Accuracy:", accuracy)
15 print("Classification Report:")
16 print(classification_rep)
```

Accuracy: 0.8489208633093526

Classification Report:

	precision	recall	f1-score	support
False	0.88	0.96	0.91	709
True	0.49	0.24	0.32	125
accuracy			0.85	834
macro avg	0.68	0.60	0.62	834
weighted avg	0.82	0.85	0.83	834

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn("The max_iter was reached which means ")



```
In [92]: ▶ 1 cv_scores_log_fss = cross_val_score(logistic_regression_model_ffs, X_
2 print('Cross-Validation Scores', cv_scores_log_fss)
3 print('Mean CV Score', cv_scores_log_fss.mean())
```

Cross-Validation Scores [0.868 0.886 0.862 0.848 0.86773547]

Mean CV Score 0.8663470941883767

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn("The max_iter was reached which means "

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn("The max_iter was reached which means "

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn("The max_iter was reached which means "

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

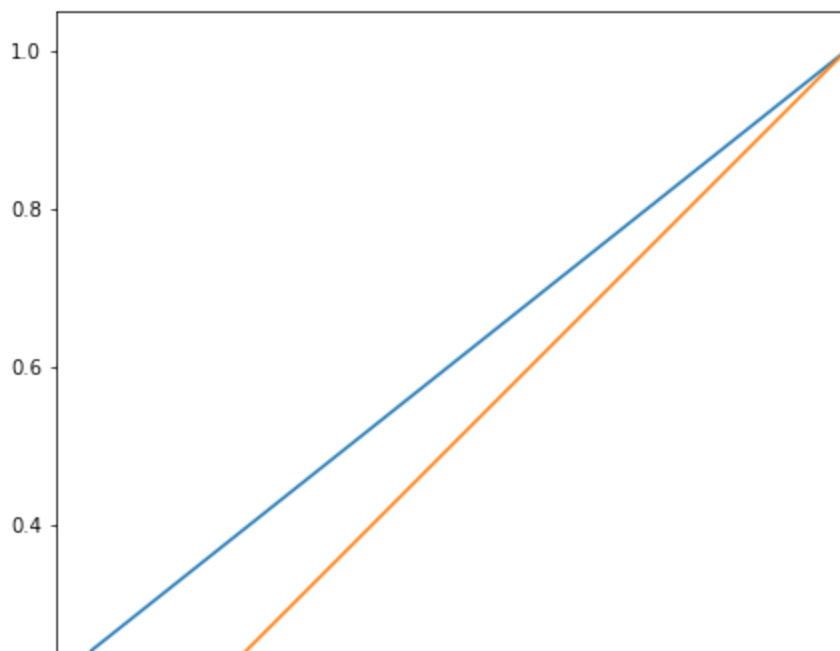
warnings.warn("The max_iter was reached which means "

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_sag.py:329: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn("The max_iter was reached which means "



```
In [93]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_df2 = logistic_regression_model_ffs.predict(my_df2_copy)
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_df2)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



Finetuning the model didn't help as much as we would have liked. Let's explore some other options.

Using Class Weighting to iterate over our model.

```

In [94]: ▶ 1 # Create the logistic regression model with class weights
2 logistic_regression_model_cw = LogisticRegression(penalty='l1', solve
3
4 # Fit the model on the resampled training data
5 logistic_regression_model_cw.fit(X_train_resampled, y_train_resampled
6
7 # Make predictions on the test data
8
9 y_pred_cw = logistic_regression_model_cw.predict_proba(X_test_final)
10 thresh = .7
11 y_pred_cw = y_pred_cw >= thresh
12 # Evaluate the model
13 accuracy = accuracy_score(y_test, (y_pred_cw[:, 1].astype(int)))
14 classification_rep = classification_report(y_test, (y_pred_cw[:, 1].a
15
16 # Print the evaluation metrics
17 print("Accuracy with Class Weights:", accuracy)
18 print("Classification Report with Class Weights:")
19 print(classification_rep)
20

```

Accuracy with Class Weights: 0.6438848920863309

Classification Report with Class Weights:

	precision	recall	f1-score	support
False	0.90	0.66	0.76	709
True	0.23	0.58	0.33	125
accuracy			0.64	834
macro avg	0.56	0.62	0.54	834
weighted avg	0.80	0.64	0.69	834



Class Weighting and Hyper Parameter Tuning performed the best on the Logistic Regression model.


```
In [95]: ► 1 logistic_regression_model_cw1 = LogisticRegression(penalty='l1', solv
2
3 # Fit the model on the training data
4 logistic_regression_model_cw1.fit(X_train_final, y_train)
5
6 # Make predictions on the test data
7 y_pred_cw1 = logistic_regression_model_cw1.predict(X_test_final)
8
9 # Evaluate the model
10 accuracy = accuracy_score(y_test, y_pred_cw1)
11 classification_rep = classification_report(y_test, y_pred_cw1)
12
13 # Print the evaluation metrics
14 print("Accuracy:", accuracy)
15 print("Classification Report:")
16 print(classification_rep)
```

Accuracy: 0.854916067146283

Classification Report:

	precision	recall	f1-score	support
False	0.90	0.93	0.92	709
True	0.52	0.41	0.46	125
accuracy			0.85	834
macro avg	0.71	0.67	0.69	834
weighted avg	0.84	0.85	0.85	834

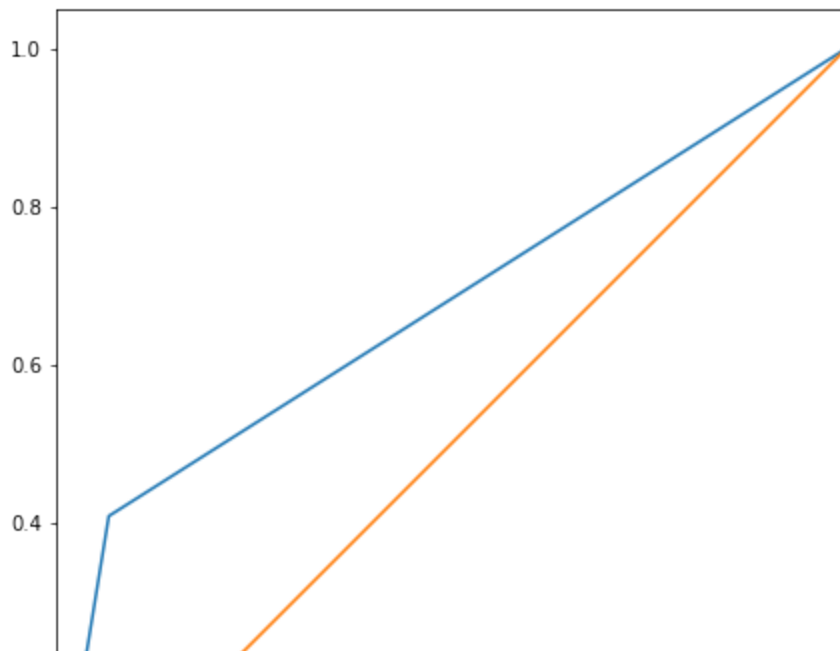


```
In [96]: ► 1 cv_scores_log_cw1 = cross_val_score(logistic_regression_model_cw1, X_
2 print('Cross-Validation Scores', cv_scores_log_cw1)
3 print('Mean CV Score', cv_scores_log_cw1.mean())
```

Cross-Validation Scores [0.866 0.874 0.834 0.844 0.86172345]

Mean CV Score 0.8559446893787575

```
In [97]: ▶ 1 # Your existing code to get the probability scores and calculate ROC
2 y_prob_df1 = logistic_regression_model_cw1.predict(my_df2_copy)
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob_df1)
4
5 # Plotting the ROC curve
6 plt.figure(figsize=(8, 6)) # Increase the figure size
7 plt.plot(fpr, tpr)
8 plt.plot([0, 1], [0, 1])
9
10 plt.xlim([0.0, 1.0])
11 plt.ylim([0.0, 1.05])
12 plt.gca().set_aspect('equal', adjustable='box') # Adjust aspect ratio
13 plt.tight_layout(pad=0) # Remove any additional whitespace
14 plt.show()
15
16 roc_auc = auc(fpr, tpr)
17 print(roc_auc)
```



Here we go. Using class weighting to address the imbalancing issue has improved our model. This improved our overall accuracy, our True Recall, and AUC score.

Final Analysis

Below we will explore the final stats on our best models. The Confusion Matrix will be added here to show another level of analysis.

Our confusion matrices will have the following format:

True Negative False Positive

False Negative True Positive

Decision Tree Classifier Confusion Matrix

In [98]: ▶

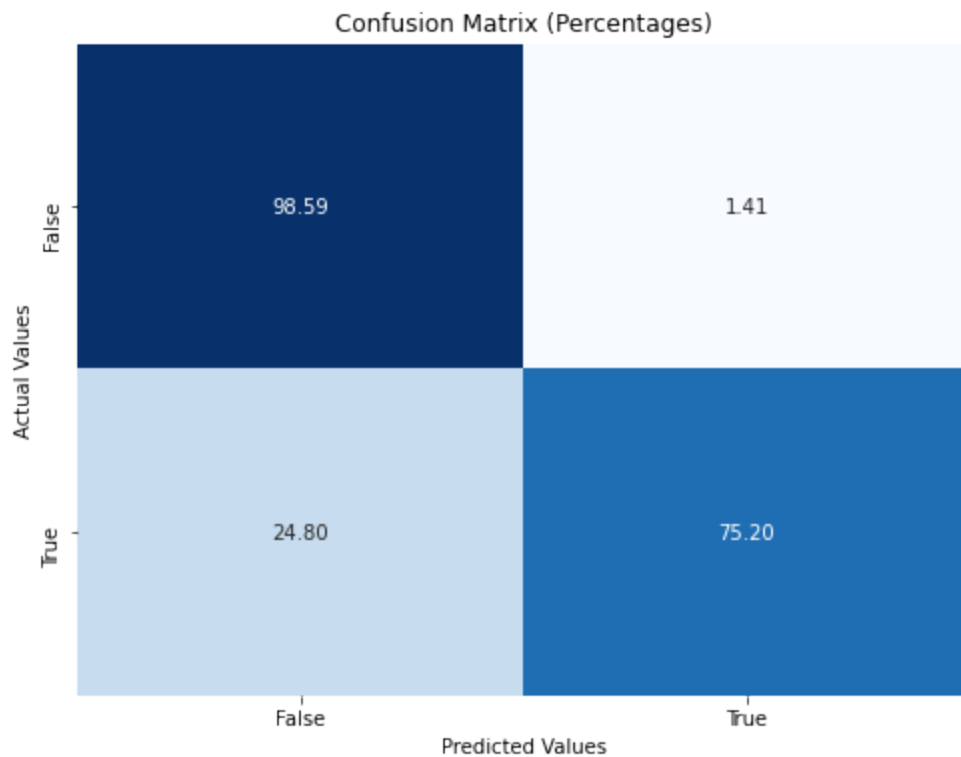
```
1 #conf_matrix_dt = confusion_matrix(y_test, y_pred_dt_fs)
2
3
4 # Define class names (modify these based on your actual class names)
5 #class_names = ['False', 'True'] # Example class names, replace with
6
7 # Plotting the confusion matrix
8 #plt.figure(figsize=(8, 6))
9 #sns.heatmap(conf_matrix_dt, annot=True, cmap='Blues', fmt='g', cbar=
10 #             xticklabels=class_names, yticklabels=class_names)
11 #plt.figure(figsize= (8,6))
12 #sns.heatmap(conf_matrix_dt, annot=True, cmap='Blues', fmt='g', cbar=
13 #plt.xlabel('Predicted Values')
14 #plt.ylabel('Actual Values')
15 #plt.title('Confusion Matrix')
16 #plt.show()
```



```

In [99]: ▶ 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from sklearn.metrics import confusion_matrix
5
6 # Assuming y_test and y_pred_dt_fs are already defined
7 conf_matrix_dt = confusion_matrix(y_test, y_pred_dt_fs)
8
9 # Normalize the confusion matrix by row (i.e., by the actual class counts)
10 conf_matrix_dt_percent = conf_matrix_dt.astype('float') / conf_matrix_dt.sum(axis=1)
11
12 # Define class names (modify these based on your actual class names)
13 class_names = ['False', 'True'] # Example class names, replace with actual
14
15 # Plotting the confusion matrix with percentages
16 plt.figure(figsize=(8, 6))
17 sns.heatmap(conf_matrix_dt_percent, annot=True, cmap='Blues', fmt='.2f',
18             xticklabels=class_names, yticklabels=class_names)
19 plt.xlabel('Predicted Values')
20 plt.ylabel('Actual Values')
21 plt.title('Confusion Matrix (Percentages)')
22 plt.show()
23

```



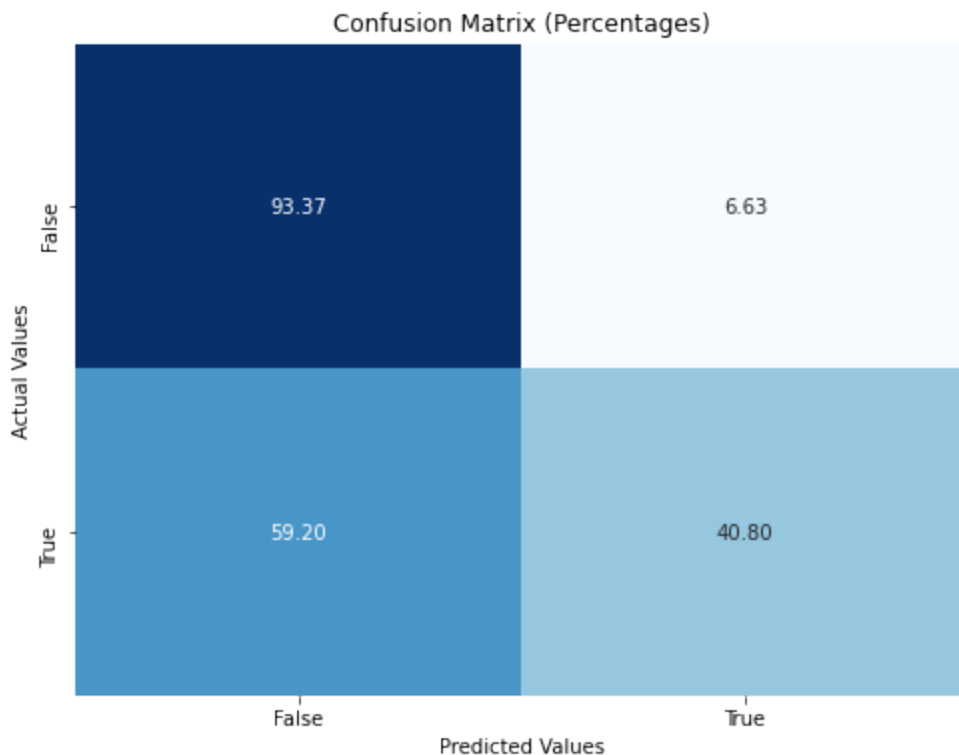
We can see the model correctly identifies 98% of True Negatives and 75% of True Positives

Below we will now take a look at our Confusion matrix for our Logistic Regression Model.

```
In [100]: ▶ 1 # Calculate the confusion matrix
2 #conf_matrix_df1 = confusion_matrix(y_test, y_prob_df1)
3
4 # Define class names (modify these based on your actual class names)
5 #class_names = ['False', 'True'] # Example class names, replace with
6
7 # Plotting the confusion matrix
8 #plt.figure(figsize=(8, 6))
9 #sns.heatmap(conf_matrix_df1, annot=True, cmap='Blues', fmt='g', cbar=
10 #             xticklabels=class_names, yticklabels=class_names)
11 #plt.xlabel('Predicted Values')
12 #plt.ylabel('Actual Values')
13 #plt.title('Confusion Matrix')
14 #plt.show()
15
```



```
In [101]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from sklearn.metrics import confusion_matrix
5
6 # Assuming y_test and y_prob_df1 are already defined
7 conf_matrix_df1 = confusion_matrix(y_test, y_prob_df1)
8
9 # Normalize the confusion matrix by row (i.e., by the actual class counts)
10 conf_matrix_df1_percent = conf_matrix_df1.astype('float') / conf_matrix_df1.sum(axis=1)
11
12 # Define class names (modify these based on your actual class names)
13 class_names = ['False', 'True'] # Example class names, replace with actual names
14
15 # Plotting the confusion matrix with percentages
16 plt.figure(figsize=(8, 6))
17 sns.heatmap(conf_matrix_df1_percent, annot=True, cmap='Blues', fmt='.1f',
18             xticklabels=class_names, yticklabels=class_names)
19 plt.xlabel('Predicted Values')
20 plt.ylabel('Actual Values')
21 plt.title('Confusion Matrix (Percentages)')
22 plt.show()
23
```



Above is our Confusion Matrix related to our Logistic Regression Model. You can see our model correctly predicted the false value 93% of the time and correctly predicted the true value 40% of the time.

Our True Positive raw number seem low, but remember that our data sets are not perfectly balanced, so there literally aren't as many opportunities for our model to correctly guess the

Summary

We have built these models with stakeholder needs in mind. We wanted to improve overall accuracy, but an important focus is also the True Positive rate of our models as this is the percent at which our model correctly predicts when a customer will churn. This will allow stakeholders to correctly anticipate customers churning and therefore understand expected revenue loss month over month.

Feature Scaling and finetuning our ccp_alpha value produced the best results for our Decision Tree. CCP_Alpha is a measure of number of nodes pruned. This is a more sophisticated method to pruning than we looked at earlier and combined with the feature scaled data it is the most powerful.

Our Decision Tree Classifier correctly predicts True values 90% of the time, had a cross validation score of 90 and AUC score of 88. Our final Decision tree improved our accuracy from 91% ->95% and more importantly it improved our True Positive rate from 71% -> 90%.

The Logistic Regression Model did not cooperate quite like we wanted and only predicted true Positives 52% of the time. Thankfully, the model still had an 85% accuracy score and cross validation value of 85. Our final Logistic Regression model improved our True Positive rate 44% -> 52% and more importantly improved our Recall for our positive cases from 9% -> 41%.

Overall, we wanted a balanced model, but when tradeoffs were to be made, they were in favor to increase true positives.



Feature Sampling proved to our best method for model improvement.

Next Steps

The Decision tree is our best model We did not utilize balancing or SMOTE, so there could be more exploration to be done in that area to find a perfect balance. I would be interested to investigate data on where the calls are being made to and from. For example, does someone who makes most of their calls within 50 miles of where they live impact churn rate? And from a modeling perspective, sampling needs to be improved